

ATLAS

Version 3.0

User Manual

Cohort definition and analytics on OMOP CDM data

Peter Hoffmann

First Edition — July 10, 2026

Preface

This is the user manual for ATLAS v3.0, the next version of the OHDSI ATLAS web application. It is a working manual: it describes what each area of the application does, how to use it, and the ATLAS-v3.0-specific behaviour that a researcher meets at the keyboard — the configuration panel, the version-history tabs, the cohort-builder controls, the plugin surface, the auto-save semantics — in the order a user encounters them.

This manual is *complementary* to *The Book of OHDSI* [20], not a replacement for it. The Book is the canonical reference for the methodology behind every ATLAS feature: the OMOP Common Data Model, the standardised vocabularies, phenotype development, characterization, comparative effectiveness, clinical and method validity. Where this manual touches one of those topics, it does so only enough to name the ATLAS controls a user is looking at, and points at the relevant Book chapter for the rest. Read them alongside each other.

Audience

The manual serves several readers:

- **Researchers and analysts** who design cohort definitions, run characterizations, pathways, or incidence rate analyses, and read the resulting reports.
- **Clinicians and domain experts** reviewing concept sets and cohort designs for clinical validity — they need enough of the interface vocabulary to comment usefully without becoming ATLAS operators.
- **Administrators** configuring data sources, managing roles, and supporting users in their organisation.
- **New OHDSI participants** for whom ATLAS v3.0 is their first contact with the OMOP Common Data Model and the OHDSI toolchain.

What this manual is, and is not

This is a *task-oriented* manual. Each chapter takes one feature, explains what it lets you do, and walks through the workflow. Where a screen requires deeper conceptual background — the difference between a standard concept and a source code, why a cohort is more than a list of people — the manual stops to explain rather than glossing over. Cross-references rather than re-explanations are used when a concept reappears in a later chapter.

This is not an installation handbook for first-time deployment; the appendix gives a working setup recipe but the canonical instructions live in the project's repository. It is

also not a WebAPI reference. The manual documents the user interface; if you need to call WebAPI directly — to script cohort generation, batch concept-set imports, or build a plugin against the same endpoints the UI uses — the appendix points to WebAPI’s live swagger UI, which is the authoritative endpoint reference. Per-section endpoint footnotes are deliberately not included.

Structure

The chapters follow the order in which most users meet the features:

- **Foundations** (Chapters 1–2) — what ATLAS v3.0 is for, and how to sign in and navigate.
- **Building blocks** (Chapters 3–5) — data sources, vocabulary search, and concept sets, the components every later workflow draws on.
- **Cohort definitions** (Chapter 6) — the longest single workflow in the application: entry events, inclusion criteria, exit, generation, reports, and the diagnostics-and-evaluation cycle the OHDSI community wraps around it.
- **Analyses** (Chapters 7–11) — characterizations, cohort pathways, incidence rates, patient profiles (including phenotype validation with PheValuator), and reusable feature analyses.
- **Operations** (Chapters 12–14) — versioning, administration, and the plugin system.
- **Appendix A** — a setup recipe for WebAPI v3.0 and the ATLAS v3.0 frontend, plus pointers to the live WebAPI reference and the wider OHDSI toolchain.

Conventions

Visible UI labels are bold (**Save**). Routes, configuration keys, and code identifiers appear in monospace (`/cohorts`). Menu paths use a sans-serif arrow form (Analysis → Characterizations). Screenshots are shown framed; until real captures replace them, slots are marked with a placeholder caption that names the figure to come. Subdued left-rule blocks distinguish notes, tips, cautions, and concept side-bars from the running prose.

A short cheat sheet of the OHDSI ecosystem (Achilles, HADES, PheValuator, Cohort-Diagnostics, Athena, Broadsea, Phenotype Library, Themis) is in the introduction’s “OHDSI toolchain at a glance” notebook. A glossary of OMOP / OHDSI terminology (cohort, concept, era, index date, standard concept, etc.) is at the back of the manual.

Acknowledgments

ATLAS v3.0 is the work of the OHDSI community continuing the long tradition of the original ATLAS. This manual draws on the OHDSI ATLAS wiki [21], *The Book of OHDSI* [20], and on the ATLAS v3.0 codebase itself. Any error in description belongs to the manual; the application owes its shape to its many contributors.

Contents

Preface	i
I Foundations	1
1 Introduction	2
1.1 What ATLAS v3.0 is for	2
1.2 What ATLAS v3.0 is not	5
1.3 How this manual is organised	5
1.4 Conventions used	5
2 Getting started	7
2.1 Before you log in	7
2.2 Signing in	7
2.2.1 Sessions and how they expire	8
2.3 The main navigation	9
2.4 Switching language	10
2.5 A first walkthrough	10
3 A first phenotype: end-to-end	11
3.1 The phenotype	11
3.2 Step 1 — Confirm the source is available	11
3.3 Step 2 — Build the metformin concept set	11
3.4 Step 3 — Build the diabetes concept set	12
3.5 Step 4 — Create the cohort definition	12
3.6 Step 5 — Add the inclusion rule	13
3.7 Step 6 — Set cohort exit	13
3.8 Step 7 — Generate	13
3.9 Step 8 — Read the inclusion-rule report	14
3.10 Step 9 — Eyeball some persons	14
3.11 Where to go from here	14
II Building Blocks	15
4 Data sources	16
4.1 Picking a source	16
4.2 Reading a source report	17

4.2.1	Dashboard	17
4.2.2	Data Density	17
4.2.3	Person	18
4.2.4	Visit	18
4.2.5	Condition / Drug / Procedure / Measurement / Observation . .	18
4.2.6	Death	19
4.3	Comparing sources	19
4.4	Data quality	19
4.5	Patient cache and cohort reports on a source	19
5	Vocabulary search	21
5.1	Where vocabulary search lives	21
5.2	Running a search	21
5.2.1	What matches	22
5.2.2	Reading the result table	22
5.3	Concept details	22
5.3.1	Header	23
5.3.2	Stat cards	23
5.3.3	Concept attributes	23
5.3.4	Hierarchy mini-map	23
5.3.5	Related concepts	23
5.3.6	Drilldown report	23
5.4	From search to set	24
5.5	Search strategies that work in practice	24
6	Concept sets	25
6.1	Why a list of concepts is not enough	25
6.2	Creating a concept set	25
6.3	The editor tabs	26
6.4	The Selected tab	26
6.5	The Included tab	27
6.6	The Source Codes tab	27
6.7	Bulk-adding concepts	27
6.7.1	Paste IDs	27
6.7.2	Import codes	27
6.7.3	Import JSON	28
6.8	Tags	28
6.9	Recommend	28
6.10	Compare	29
6.11	Saving	30
6.12	Starting from a phenotype library	30
6.13	Discovery	31

III Cohort Definitions	32
7 Cohort definitions	33
7.1 Live preview	33
7.2 Listing and creating cohorts	35
7.2.1 Duplicating a cohort	35
7.3 The four logical pieces of a definition	35
7.4 Entry event (primary criteria)	35
7.4.1 Cohort entry observation window	36
7.4.2 Three qualifying-event limits	37
7.5 Inclusion criteria	37
7.5.1 Rail and detail layout	37
7.5.2 Anatomy of a rule	38
7.5.3 Grouping criteria with AND / OR	38
7.5.4 The criteria event card	39
7.5.5 Reading the inclusion report	40
7.6 Cohort exit and censoring	41
7.6.1 Exit criteria	41
7.6.2 Censoring events	41
7.6.3 Era collapse	41
7.7 Concept sets used by this cohort	42
7.8 Generation	43
7.8.1 Cohort samples	43
7.9 Going further: external diagnostic tools	44
7.10 Saving, copying and exporting	45
7.10.1 Tagging a cohort	45
7.10.2 Auto-save and draft recovery	45
IV Analyses	47
8 The Analysis Hub	48
8.1 Where the analyses live	48
8.2 Lists, editors, and results	48
8.3 A note on Profiles	49
9 Characterizations	50
9.1 What you can answer with a characterization	50
9.2 Anatomy of a characterization	51
9.2.1 Building a subgroup	51
9.3 Generating and viewing results	51
9.3.1 Result view modes	52
9.4 Export	53
9.5 Common pitfalls	53

9.6 Reproducing a characterization in R	55
10 Cohort pathways	56
10.1 When this analysis is the right tool	56
10.2 Setting up a pathway	56
10.3 Reading the result	57
11 Incidence rates	59
11.1 Setting up an incidence rate analysis	59
11.2 Reading the result	60
11.2.1 The Insights rail	61
12 Profiles	63
12.1 Finding a person	63
12.2 The timeline view	63
12.3 The rest of the profile page	64
12.3.1 The Demographics strip	64
12.3.2 The Observation periods band	64
12.3.3 The Highlights panel	65
12.4 Validating a phenotype using Profiles	65
12.5 Going further: PheValuator	67
13 Feature analyses	68
13.1 Built-in versus custom features	68
13.2 The three feature-analysis types	68
13.3 Reuse	69
V Operations	70
14 Versions	71
14.1 Where versions live	71
14.2 Previewing a version	72
14.3 Restoring a version	72
14.4 Pinning a published study	72
15 Administration	73
15.1 The configuration panel	73
15.2 Caches and TrexSQL	74
15.3 Roles and permissions	75
15.3.1 Built-in roles	75
15.3.2 Editing a role	76
15.3.3 Creating, importing and deleting a role	76
15.3.4 Permission reference	76
15.4 Jobs and long-running tasks	77

15.5 Audit information	78
16 Plugins	79
16.1 What plugins look like	79
16.2 Common plugin categories	79
16.3 Trust and isolation	79
16.4 The <code>plugins.json</code> manifest	80
16.5 Errors and troubleshooting	82
17 Troubleshooting	83
17.1 I cannot log in	83
17.2 I cannot see any data	84
17.3 Cohort generation does not produce what I expect	84
17.4 Concept set issues	85
17.5 Permissions: I cannot see or do something	85
17.6 Errors that show up in the browser console	86
17.7 Plugin problems	86
17.8 When the answer is not in this manual	86
A Setup guide: WebAPI and ATLAS v3.0	87
A.1 The pieces	87
A.2 Prerequisites	88
A.3 The Docker Compose path (recommended for evaluation)	88
A.4 Native install	89
A.4.1 1. Database	89
A.4.2 2. WebAPI v3.0	89
A.4.3 2a. Enable the TrexSQL query cache	90
A.4.4 3. Add data sources to WebAPI	91
A.4.5 4. ATLAS v3.0 frontend	92
A.4.6 Externalized runtime configuration	92
A.5 First-time login and creating a user	93
A.6 Common installation issues	93
A.7 Verifying a clean setup	94
A.8 Achilles: the missing prerequisite for descriptive reports	94
A.9 Related: Broadsea	95
A.10 Further reading and API reference	95
Glossary	96
Bibliography	98

PART I

Foundations

Introduction

ATLAS v3.0 is a web application for designing, generating, and reviewing patient cohorts on data that conforms to the OMOP Common Data Model (CDM) [4, 16, 29]. It is a modern reimplementaion of the long-standing OHDSI ATLAS tool. Functionally it covers the same ground as ATLAS 2.x — vocabulary search, concept sets, cohort definitions, characterizations, cohort pathways, incidence rates and patient profiles — built on top of the same WebAPI backend. The user interface has been rebuilt with Vue 3 [43] and Vuetify [8], and several workflows have been streamlined.

This manual is task-oriented. Each chapter takes one feature, explains what it lets you do, and walks through it in the order a user would meet it. If you are completely new to OHDSI, read this introduction and Chapter 2 first. If you already know ATLAS, you can jump straight to the chapter that matches the feature you need.




Figure 1.1. The ATLAS v3.0 home page. The top navigation lists the major application areas; the home card provides the two most-used entry points (**Define a New Cohort**, **Search the Vocabulary**) plus a row of analysis-area shortcuts.

1.1 What ATLAS v3.0 is for

Observational health data — claims, electronic health records, registries — answer different questions than clinical trials. They are large, messy, and the records were collected


for billing or care, not research. The OHDSI community addresses this with two pieces: a standardised data model (the OMOP CDM) and a library of methods that operate on that model. ATLAS v3.0 sits on top of both. It is the place where a researcher:

- explores which data sources are available and what they contain;
- searches the standardised vocabularies for the concepts that describe a clinical idea (a disease, a drug class, a procedure);
- collects those concepts into reusable *concept sets*;
- composes *cohort definitions* — formal, executable descriptions of a population such as "new users of metformin without a prior diagnosis of diabetes";
- generates the cohort against one or more data sources and reviews descriptive reports about it;
- runs additional analyses — characterizations, pathways, incidence rates — using cohorts as inputs.

Everything ATLAS v3.0 produces is a portable artefact. A cohort definition is stored as JSON and SQL and can be shared, version-controlled, and run again on a different network on a different CDM database to reproduce a finding. The OHDSI design ethos is to *execute analytics in place*: the patient-level data never leaves the hosting institution, and only aggregate results are pooled across the network [20].

The cohort builder in ATLAS v3.0 is *rule-based*: a person is in the cohort if and only if they satisfy a set of explicit, hand-crafted criteria. Probabilistic phenotyping — where a model assigns each person a probability of meeting the phenotype — is a separate practice supported by the HADES PheValuator package (Chapter 12) [37].

The OMOP CDM is a person-centric relational schema. Every record about every person — a condition, a drug exposure, a measurement result — is stored in a small, predictable set of tables, and every clinical idea is represented by a *concept* drawn from the OHDSI standardised vocabularies. Because the schema and the vocabularies are the same everywhere, the same cohort definition you build in ATLAS v3.0 produces comparable results on a Korean claims database, a Dutch GP record system, and a US academic medical centre [41]. See *The Book of OHDSI*, chapter "OHDSI Analytics Tools" [20] for the longer treatment.



How ATLAS v3.0 sits in the OHDSI ecosystem. Solid arrows are runtime data flows. Two further relationships are not drawn to keep the figure readable: ATLAS reads Achilles' summary tables when it renders descriptive reports, and ATLAS imports cohort JSON from the Phenotype Library.

Note

The OHDSI toolchain at a glance. ATLAS is one piece of a larger OHDSI ecosystem. The names below are all referenced elsewhere in this manual; this is the cheat sheet:

ATLAS this web application; designs and runs cohort definitions and analyses.

WebAPI the Java/Spring backend ATLAS talks to. It compiles a cohort definition into SQL and executes it against a CDM source.

OMOP CDM the standardised relational schema all CDM sources conform to.

Athena the OHDSI vocabulary repository. The standardised vocabularies are released here.

Achilles an R package that profiles a CDM source and writes summary tables. ATLAS's descriptive reports read from those tables — if Achilles has not been run, the Data Sources page is blank.

HADES the OHDSI suite of R analytical packages (CohortMethod, FeatureExtraction, PatientLevelPrediction, and the two below).

CohortDiagnostics HADES package that runs deeper diagnostics on a cohort definition exported from ATLAS.

PheValuator HADES package that estimates phenotype sensitivity, specificity, and PPV against a probabilistic gold standard.

OHDSI Phenotype Library community-curated repository of vetted cohort definitions, each release DOI'd.

Themis the OHDSI conventions group; defines best practices for CDM and cohort design.

Broadsea a Docker bundle of the OHDSI stack (currently ATLAS 2.x; v3.0 support pending).

1.2 What ATLAS v3.0 is not

ATLAS v3.0 is a *design and review* tool. It does not store the patient data itself, and it does not run advanced analytics in the browser. When you press **Generate**, ATLAS v3.0 sends the cohort definition to WebAPI, which compiles it to SQL and runs it against the chosen CDM database.

The Population-Level Estimation and Patient-Level Prediction [33] modules from older ATLAS versions are *not* part of ATLAS v3.0 and there is no concrete plan to re-implement them in this codebase. For comparative effectiveness and prediction work, use the OHDSI HADES Libraries [23] directly in R against the same CDM data, or fall back to the classic ATLAS instance your organisation maintains.

1.3 How this manual is organised

- **Chapter 2, Getting Started** — what you need before you log in, how to log in, and a tour of the main navigation.
- **Chapter 4, Data Sources** — picking a CDM database and reading its descriptive reports.
- **Chapter 5, Vocabulary Search** — searching the OMOP standardised vocabularies.
- **Chapter 6, Concept Sets** — building and sharing reusable collections of concepts.
- **Chapter 7, Cohort Definitions** — the longest chapter; covers entry events, inclusion rules, exit, censoring, generation, and reports.
- **Chapters 9–11** — the analytical modules: characterizations, pathways, incidence rates.
- **Chapter 12, Profiles** — exploring an individual patient timeline.
- **Chapter 13, Feature Analyses** — reusable covariate definitions consumed by characterizations.
- **Chapter 14, Versions** — saving, previewing and restoring previous versions of cohorts and concept sets.
- **Chapter 15, Administration** — roles and permissions for users with administrative access.
- **Chapter 16, Plugins** — what plugins are and how they appear inside ATLAS v3.0.

1.4 Conventions used

- UI labels are bold: click **Save**, open the **Concept Sets** tab.
- Menu paths use chevrons: **Analysis** » **Characterizations**.
- Routes and code are monospaced: `/cohorts`, the JSON key `conceptSet`.
- Side notes appear in coloured boxes:

Note

A short fact worth knowing but not essential to the procedure being described.

A faster way to do the same thing.

Something that can ruin your work or produce misleading results if you are not careful.

Throughout the manual, where a section maps onto a single WebAPI endpoint or store action, that information appears as a small footnote at the end of the section so power users know where the UI's data is coming from.

Getting started

This chapter assumes someone has already installed ATLAS v3.0 and pointed it at a WebAPI instance — installation is covered in the project’s repository README. Here we look at the experience from a new user’s perspective: what you need on day one, how to log in, and what each main area of the application is for.

2.1 Before you log in

You need three things:

1. A URL for the ATLAS v3.0 instance your organisation hosts. In Docker development this is usually `https://localhost`.
2. Credentials, or an identity provider (LDAP, OAuth, SAML, OpenID Connect) that the WebAPI is configured to accept.
3. At least one CDM database that WebAPI is connected to — without a data source, every analytical feature is read-only and most lists will be empty.

Note

Public OHDSI demonstration instances such as <https://atlas-demo.ohdsi.org> expose synthetic data. They are useful for trying ATLAS v3.0 risk-free but the population sizes there are not representative of the data you would use for a real study.

2.2 Signing in

Open the ATLAS v3.0 URL in a modern browser (Chromium-based browsers and Firefox are the test targets). At startup the app always calls `user/me`. When the WebAPI is configured for *anonymous access* that call returns an anonymous permission set, no login prompt appears, and you can browse whatever the anonymous role is allowed to see — the server still gates every action per-endpoint. The shipped Docker deployment instead turns authentication on through its runtime config (`userAuthenticationEnabled: true`), so a session is required before you act on WebAPI data.

When authentication is on, the login is a *modal dialog* rather than a separate page: the router guard opens it the first time you hit a gated action, for example clicking into **Cohort Definitions**.

What appears in the login dialog depends on what the WebAPI is configured to accept:




Figure 2.1. Login dialog. Provider buttons appear when the deployment is configured for OAuth, SAML, or OpenID; here only the database-credentials form is enabled.

Username and password A simple form. Used by LDAP-backed deployments and by the local user database.

OAuth providers One button per configured provider, e.g. Google. Clicking opens the provider's consent screen and returns to ATLAS v3.0.

SAML / OpenID Connect Concrete supported single-sign-on flows. The provider redirects back to a dedicated callback route (`/saml/callback` or `/openid/callback`) which finishes the sign-in and drops you back where you started.

After a successful sign-in your username appears in the top right. Click it to sign out.

Note

Deep links. ATLAS v3.0 honours two query parameters when the app loads: `?cohortId=X` opens that cohort directly, and `?route=/path` navigates to an arbitrary in-app route. They are processed once, on the first navigation after boot — handy for sharing a link straight to a specific artefact.

2.2.1 Sessions and how they expire

ATLAS v3.0 stores a JWT bearer token in the browser. The token has a fixed lifetime; a few minutes before it expires ATLAS v3.0 shows a dialog inviting you to extend the session. If you walk away from your machine the token quietly expires and the next gated action re-opens the login modal (or, on a deployment configured for anonymous access, silently drops you back to the anonymous permission set). No unsaved work is

lost — ATLAS v3.0 keeps an in-progress cohort definition in session storage and restores it after you sign back in.

If you keep getting logged out unexpectedly, check the system clock on the WebAPI server. A clock skew larger than the token's clock-tolerance makes every freshly issued token look already expired.

2.3 The main navigation

The top of the screen has a navigation bar with the major application areas, plus a configuration menu for users with administrative access. The left sidebar (where present) gives a chapter-by-chapter breakdown of the area you are currently in.



Figure 2.2. The top navigation bar. From left: the ATLAS logo, the five built-in entries, then the right-hand cluster of Feedback button, language selector, configuration cog (visible to admins only), and the user menu.

The top navigation bar contains five built-in entries plus a small control cluster on the right. From left to right:

Data Sources The list of connected CDM databases, with a descriptive report per source. Covered in Chapter 4.

Concept Sets Browse, search and edit reusable collections of concepts. Vocabulary search lives inside the same area, on a sibling tab. Covered in Chapter 5 and Chapter 6.

Cohort Definitions List, search, filter, edit and generate cohorts. Covered in Chapter 7.

Analysis A tabbed area combining Analysis > Characterizations, Analysis > Cohort Pathways, Analysis > Incidence Rates and Analysis > Feature Analyses. Each is its own chapter (9–13).

Profiles Per-patient timeline view. Reached either directly from the top nav, or via a "view profile" action from another page. Covered in Chapter 12.

The right-hand control cluster carries a **Feedback** button (opens an external feedback form configured by the deployment), a language selector, the configuration cog, and the user menu. The cog is only shown when the signed-in user holds at least one administrative permission — if you cannot see it, ask your administrator. Plugins your organisation has installed extend this navigation with additional entries (Chapter 16).

Note

Permissions affect what you see. Many actions in ATLAS v3.0 are gated by per-entity permissions: you may see a cohort or concept set in a list but find that **Save** or **Delete** is disabled, and the configuration cog is hidden entirely if you do not have any admin permission. Chapter 15 describes the role and permission model.

A few features live outside the top nav and are reached by other means:

Home Click the ATLAS logo on the far left to return to the landing page at any time.

Jobs Long-running task status lives inside the configuration panel (cog icon) under the Jobs tab.

2.4 Switching language

ATLAS v3.0 ships its translations *bundled* with the frontend — the JSON files under `src/locales/` (English, Korean, Russian, Chinese and Japanese). They are not fetched from WebAPI at runtime. Pick a language from the language selector in the header control cluster — the same control described under *The main navigation*, not the user menu; the interface re-renders without reloading the page. Keys a given locale does not translate fall back to bundled English. The text of cohort definitions and concept sets is stored verbatim — only the chrome around them is translated.

2.5 A first walkthrough

To get a feel for the application, the OHDSI community recommends this short loop on a demo dataset:

1. Open **Data Sources** and pick any source. Wait for the descriptive report to load and skim the demographics chart.
2. Open **Concept Sets**, switch to the **Concept Search** tab and search the vocabulary for a familiar condition (e.g. "atrial fibrillation"), add a couple of concepts to a new set, tick **Descendants** on each row in the **Included Concepts** tab and save.
3. Open **Cohort Definitions**, click **New Cohort**, set your concept set as the entry event, leave inclusion criteria empty, save, then expand the **Generation** section and click **Run** on the same data source's row.
4. When the row reaches **COMPLETED**, click its **Inclusion report** action — the report drawer shows the attrition funnel and per-rule breakdown.

That round-trip — vocabulary → concept set → cohort → report — is the core motion of ATLAS v3.0. The rest of the manual is variations on it.

A first phenotype: end-to-end

This chapter walks through one complete piece of work in ATLAS v3.0, from logging in to reading the cohort report. The reference chapters that follow describe each feature in depth; this one shows how they connect. If you are new to ATLAS, do this walkthrough on the public demo instance (<https://atlas-demo.ohdsi.org>) before tackling your own study.

Note

The demo instance ships a synthetic source called SYNPUF1K that is small enough to generate against in seconds. The exact counts you see may differ slightly from the figures below if the demo instance refreshes its sample, but the *shape* of the result will be the same.

3.1 The phenotype

Goal: *new users of metformin who do not have a prior diagnosis of diabetes recorded in the year before they start the drug.* This is a deliberately simple definition that touches every part of ATLAS: vocabulary search, two concept sets, an entry event with a washout, one inclusion rule using a zero-cardinality "absence" pattern, and a generation against the demo source.

3.2 Step 1 — Confirm the source is available

Open **Data Sources** from the top navigation. Pick SYNPUF1K from the source dropdown; the **Dashboard** report should render demographics and an observation-period chart. If the dashboard is empty, see Chapter 17.

3.3 Step 2 — Build the metformin concept set

1. Open **Concept Sets** from the top navigation; the **Concept Sets** list-tab is the default. Click **New Concept Set**; the editor opens as a side drawer. Name it `Metformin --- ingredient` in the inline title at the top.
2. In the editor drawer, switch to the **Search** tab. Type `metformin` into the search hero and press Enter.

3. In the results table, sort or scan to find the row whose name is exactly *metformin*, vocabulary RxNorm, type *Standard*, class *Ingredient* (concept ID 1503297). Click **Add** on that row.
4. Switch back to the **Included Concepts** tab. Tick **Descendants** on the row. Leave **Mapped** off and **Exclude** off.
5. Skim the row list — with descendants the set should cover branded and generic metformin products.
6. Click **Create** (or **Save**) on the drawer header to persist the set.

3.4 Step 3 — Build the diabetes concept set

The same flow with a different parent concept:

1. Click **New Concept Set**, name it *Diabetes mellitus*.
2. On the **Search** tab, search *diabetes mellitus*; find the SNOMED standard concept whose name is exactly *Diabetes mellitus* (concept ID 201820). Click **Add**.
3. Switch to the **Included Concepts** tab and tick **Descendants** on the row so all subtypes are covered.
4. Click **Create**.

3.5 Step 4 — Create the cohort definition

1. Open **Cohort Definitions** from the top navigation. Click **New Cohort**, name it *New users of metformin without prior diabetes*, and save.
2. In step 1 of the editor (**Cohort Entry Events**), click **Add criteria to group...** and pick **Drug Exposure** from the menu.
3. In the entry-event card, click the concept-set picker and choose *Metformin --- ingredient*.
4. Click the observation-period chip ("0 days before and 0 days after the event index date") to open its dialog; set the prior days to 365, leave post days at 0, and close the dialog.
5. Set the **Cohort entry on** toggle (top-right of step 1) to **earliest event** so that only each person's first qualifying drug exposure counts.

3.6 Step 5 — Add the inclusion rule

If the demo source's patient cache is built, expand the **Generation** section at the bottom of the cohort builder and pick SYNPUF1K in **Live preview cache**. The inclusion rail then shows running counts per rule as you edit — the "no prior diabetes" rule should drop the count noticeably when applied. If it does not, either the source has very few diabetes records or the concept set has not resolved correctly; treat the live counts as a sanity check.

1. In step 2 (**Inclusion Criteria**), click **Add rule**. Name the new rule No prior diabetes diagnosis.
2. Inside the rule, click **Add criteria to group...** and pick **Condition Occurrence** from the menu.
3. For the concept set, pick Diabetes mellitus.
4. In the temporal window, choose "365 days before index" to "0 days before index".
5. In the cardinality dropdown, pick *exactly 0 occurrences*. This is the no-negation pattern from Chapter 7: there is no NOT operator; absence is expressed as a zero-cardinality rule.

3.7 Step 6 — Set cohort exit

In the **Cohort Exit & Eras** section, set the Strategy toggle to **Drug exposure** (tooltip: "End of continuous drug exposure") so the person is in the cohort for as long as they keep getting metformin prescriptions (allowing a 30-day persistence window between consecutive fills). Pick Metformin --- ingredient as the concept set.

3.8 Step 7 — Generate

Click **Save** to persist the cohort, then expand the **Generation** section at the bottom of the builder. In the data-source row for SYNPUF1K, click **Run**; the row's status chip flips to **RUNNING** and updates as the job progresses. On the demo source the run completes in seconds and the row shows **COMPLETED** with a patient count.

Note

If the run finishes but the patient count is blank, or the row sits at **RUNNING** forever, the troubleshooting chapter (Chapter 17) covers the usual culprits: source not Achilles-processed, a patient cache that needs to be built first, no rows because the concept set resolves to nothing on this source, or an attribute filter that excludes everyone.

3.9 Step 8 — Read the inclusion-rule report

When the SYNPUF1K row shows **COMPLETED**, click its **Inclusion report** action. The report drawer slides in from the right and shows the attrition funnel together with a per-rule satisfaction table:

- how many persons in the source had any qualifying drug-exposure entry event,
- how many of those passed the prior-observation window, and
- how many remained after the "no prior diabetes" rule was applied.

If any rule drops your numerator to zero, that is the rule to investigate first. If the diabetes-exclusion rule drops nobody, either there are no diabetes records before metformin in this source (unlikely) or the concept set resolved to nothing (check it in step 3).

3.10 Step 9 — Eyeball some persons

Close the inclusion-report drawer and click **Samples** on the same SYNPUF1K row. The samples panel opens in the drawer; click **New sample** to draw a random subset of the cohort, then click any person to open their profile (Chapter 12) inside the same drawer, with the membership window shaded on the timeline. The drawer's back-arrow returns to the samples list. Skim a few patients: do they look like new metformin users without a prior diabetes record?

3.11 Where to go from here

Variations on this same flow cover most of what ATLAS does:

- Add a second target cohort (e.g. new users of sulfonylureas) and run a *characterization* across both — Chapter 9.
- Define an outcome cohort (e.g. acute kidney injury) and run an *incidence rate* analysis with metformin as the target — Chapter 11.
- Define a small set of event cohorts (insulin, sulfonylurea, GLP-1 agonist, SGLT2 inhibitor) and run a *cohort pathway* analysis from the metformin target — Chapter 10.
- Open one of the cohort's persons in **Profiles** to see the actual records that triggered their cohort entry — Chapter 12.

The reference chapters that follow document each of these features in depth.

PART II

Building Blocks

Data sources

The **Data Sources** area lists every CDM database that the WebAPI is connected to and gives a descriptive view into each. It is the right place to start when you join a new project: before designing a cohort, you want to know what is in the data — how many persons, what condition mix, what time span, and how complete each domain is.

4.1 Picking a source

Open **Data Sources** from the top navigation. The page has two permanent controls:

- a *source picker* dropdown in the page header, showing every CDM database the WebAPI is connected to;
- a left-rail *report sidebar* listing the descriptive reports available for the selected source.




Figure 4.1. Data Sources page. The source-picker dropdown sits in the top-right; the left rail lists report types grouped by Overview, Demographics and Clinical Domains; the main panel renders the selected report.

The dropdown is populated from the WebAPI `/source/sources` endpoint; a source you cannot see may exist but be hidden from your role. Until a source is picked the sidebar is disabled and the main panel shows a "select a data source" hint.

Note

The "source key" — the short text identifier such as CCAE or SYNPUF1K — appears throughout the URL bar and the JSON of generated cohorts. When two analysts compare results, agreeing on the source key is the unambiguous way to know they are on the same database.

Descriptive reports on a source require that the OHDSI Achilles [9] routine has run against that source's results schema. A source that is configured in WebAPI but has not been Achilles-processed shows up in the list but renders empty report tabs. If your data source page looks bare on every tab, ask your administrator whether Achilles has been run — this is the single most common "set up but nothing shows" cause.

4.2 Reading a source report

Each source exposes the same family of descriptive reports, picked from the left sidebar. The thirteen built-in entries are grouped as:

Overview Dashboard, Data Density.

Demographics Person, Observation Period, Death.

Clinical Domains Visit, Condition Occurrence, Condition Era, Procedure, Drug Exposure, Drug Era, Measurement, Observation.

A report that has no Achilles results behind it on the chosen source still opens, but renders empty.

4.2.1 Dashboard

The default tab. It summarises the source at a glance: total persons, gender distribution, year-of-birth distribution, observation period density, and a cumulative observation chart. Use it as a sanity check. A source with millions of persons but only a couple of years of follow-up makes a different kind of study possible than a source with hundreds of thousands of persons followed for two decades.

Large counts throughout the dashboard charts — in both tooltips and axis labels — are rendered in compact SI form: 1,000 shows as 1.0k, 3,000,000 as 3.0M, scaling up to T, so a multi-million-person source stays readable at a glance. The cumulative observation curve is a percent-of-persons metric and scales to a full 100% on its y-axis. The age and year-of-birth histograms render cleanly with the same compact formatting.

4.2.2 Data Density

A breakdown of how many records exist per domain (condition, drug exposure, procedure, measurement, observation, visit, death) and how that volume changes over calendar time.




Figure 4.2. Source dashboard: CDM summary, gender split, age at first observation, and the cumulative observation curve.

Sparse rows here usually indicate either a data quality issue or that the source genuinely does not capture that domain.

4.2.3 Person

Demographics in more detail: age at first observation, race, ethnicity, where these are recorded.

4.2.4 Visit

Visit type and visit length distributions. A claims database is dominated by outpatient visits; an inpatient EHR will look very different here.

4.2.5 Condition / Drug / Procedure / Measurement / Observation

One tab per clinical domain, each showing the most prevalent concepts in that domain along with a per-concept drilldown. Scrolling these tabs is the fastest way to learn the "shape" of a source — for example, whether oncology is well represented, whether labs are coded with LOINC, whether inpatient procedures dominate over outpatient.

The default view is a *treemap*: a rectangle filled with nested boxes whose area is proportional to record count, coloured by parent vocabulary class. Hover any tile for the concept name and exact count; click into it to descend into the hierarchy. A toggle next to the treemap switches to a sortable, searchable table view with a concept-name filter — useful when you know the term you are looking for and want to find it directly rather than spotting it visually.

Click any concept row (in either view) to drill into a per-concept report: temporal trend (how the concept's frequency moves over calendar years), gender breakdown, and

the distribution of age at first occurrence. The drill-down is the right view when you suspect a concept's counts changed because of a coding-system rollover rather than a true epidemiological shift.

4.2.6 Death

Counts of recorded deaths and their cause-of-death distribution where available. Many ambulatory data sources do not capture death well; that is something to know before designing a survival study.

4.3 Comparing sources

When you have access to several sources you can run the same analysis across all of them; the comparison is what gives observational evidence its strength. From the source list, your usual workflow is:

1. open the dashboard for each candidate source;
2. note the rough size and time span;
3. open the domain tab most relevant to your study (**Drug** for a drug exposure cohort, **Condition** for a disease cohort);
4. check that the concepts you plan to use actually appear with non-trivial counts.

If the prevalence of an entry concept is zero on a source, no cohort generation will ever produce non-empty rows there — even if the SQL runs without error. Catch this in the source report rather than after a long generation.

4.4 Data quality

ATLAS v3.0 itself is not a data-quality tool, but the same Achilles results that drive its descriptive reports also flag the most common data-density anomalies: implausible dates, sudden coding-rate changes, records outside observation periods, and other signals that something upstream of the CDM has shifted. Watch for sharp discontinuities on the temporal trend charts; they usually indicate an ETL change rather than a true epidemiological event.

For a deeper, systematic view, OHDSI ships the *Data Quality Dashboard* (DQD) [14] as a sister tool to ATLAS. DQD runs over 1,500 checks organised by the Kahn data-quality framework [7] (Conformance, Completeness, Plausibility) and produces a per-source dashboard. DQD is not part of ATLAS v3.0; it runs in R against the same CDM and writes its output where your administrator hosts it.

4.5 Patient cache and cohort reports on a source

Some descriptive reports and per-source cohort reports rely on a TrexSQL patient cache being built in advance. Cache management is not exposed on the source page itself;

administrators trigger and monitor cache rebuilds from the configuration panel's **Cache** tab (Chapter 15). If a report depends on a missing cache, the report panel itself surfaces the appropriate hint.

The same source detail page is reused later when you open a generated cohort and ask for its reports — see Chapter 7.

Vocabulary search

Every clinical idea expressed in ATLAS v3.0 — a disease, a drug, a procedure — is represented by one or more *concepts* drawn from the OHDSI standardised vocabularies. Vocabulary search is the way you find those concepts. It is also the entry point to building concept sets (Chapter 6).

5.1 Where vocabulary search lives

ATLAS v3.0 surfaces vocabulary search on the **Concepts** page, which has two tabs: **Concept Sets** (the default — a list of saved concept sets) and **Concept Search**. The Concept Search tab is the standalone search experience; the same search is also embedded inline inside the concept-set editor and wherever else a workflow asks you to pick a concept (for example, when you set an entry event in the cohort builder).




Figure 5.1. Concepts page on the **Concept Sets** tab. The sibling **Search** tab hosts the standalone vocabulary search.

5.2 Running a search

Switch to the **Concept Search** tab. Type a free-text query — a name, a synonym, a code — of at least three characters into the search hero and press **Enter**. The results table fills with matching concepts.

5.2.1 What matches

The search runs against concept names, synonyms, and source codes (ICD-9, ICD-10, SNOMED, RxNorm, LOINC, NDC, and so on, depending on which vocabularies your WebAPI has loaded). A query like `atrial fibrillation` matches the standard SNOMED concept and its descendants; a query like `I48` matches the ICD-10 code and the standard concept it maps to.

Note

The OHDSI vocabulary lifecycle — standard vs. non-standard concepts, what each **Invalid Reason** flag (D, U, R) means, the **Maps to** and **Maps to value** relationships, the role of vocabularies like SNOMED, ICD-10, RxNorm and LOINC — is documented in *The Book of OHDSI*, chapter 5 "Standardized Vocabularies" [20]. This chapter only covers what the ATLAS UI shows.

5.2.2 Reading the result table

The result table has these columns: ID, Code, Name, Vocabulary, Type (the standard-concept flag), Domain, Class, Validity, and four record-count columns abbreviated **RC** (record count for the concept itself), **DRC** (descendant record count), **PC** (person count), and **DPC** (descendant person count). All columns are sortable.

To narrow a noisy search, a **Filters** menu button sits above the table; its badge shows the number of active filters. Opening it reveals multi-select facets — Vocabulary, Domain, Standard, Class, and Validity — applied client-side to the results already loaded. Each active facet value appears as a removable chip below the button, and a **Clear all** action drops them in one click. Combine the facets with column sorting — often on **DRC** to surface the most-used concepts in the source — to home in on the concept you want.

The record-count columns reflect the data source the page is configured against. If they remain blank, the source has no patient cache built (Chapter 15) or you do not have permission to read it.

The concept-name cell in every row is a link. Clicking it opens the **Concept Details** panel for that concept (next section); the same is true wherever ATLAS v3.0 lists concepts — the **Selected** tab of a concept set, the **Concept Sets** list page, the popup picker inside the cohort builder.

5.3 Concept details

Clicking a concept-name link opens **Concept Details** as a slide-in drawer covering most of the viewport. The same panel is also available as a full-page route at `/concept/<source>/<id>`, which is what you copy-paste when you want to share a concept with a colleague.

5.3.1 Header

The sticky header carries the concept name plus chips for the standard-concept flag (**Standard**, **Classification**, or **Non-standard**), vocabulary, concept ID, domain and concept class. Two action buttons sit on the right:

Add to Concept Set adds the concept to the currently-edited concept set, the same operation as the per-row **Add** action in the search table. It is always enabled: if no concept set is open, ATLAS auto-creates a new untitled set, opens the editor, and shows a snackbar reading "Added '...' → New concept set".

Copy ID writes the concept ID to the clipboard.

5.3.2 Stat cards

Below the header, one stat card per configured data source shows record count, distinct person count, and descendant record count (formatted with K / M suffixes). This is the cross-source view the search-result table cannot give you: the same concept can be common in one source and absent in another.

5.3.3 Concept attributes

The **Concept Details** card lists the full attribute set — concept ID, code, domain, vocabulary, concept class, standard concept flag, invalid reason. This is the canonical place to read off a concept's metadata; the attributes that fit in the search-results table are a subset of this list.

5.3.4 Hierarchy mini-map

The hierarchy card shows up to three parents and six children of the concept, drawn as a tree. **View full** → opens the parent-and-descendant view in a wider context for concepts with deep hierarchies.

5.3.5 Related concepts

The **Related concepts** table lists relationship rows (*Maps to*, *Subsumes*, *Has answer*, etc.) with the related concept's ID, name, vocabulary and standard flag. Clicking a related concept name opens its details in the same drawer, replacing the current concept; the back-arrow returns to the previous one. This is how you walk a relationship graph without leaving the panel.

5.3.6 Drilldown report

The **Drilldown Report** card shows the concept's prevalence on a chosen data source. A source dropdown sits at the top of the card; two tabs underneath select what to plot:

Age at first occurrence a boxplot of the age distribution at the patient's first record of the concept.

Calendar month a line chart of prevalence per 1000 people by calendar month.

The card hides itself for concepts in domains with no drilldown data (e.g. concept-class concepts that never appear as records).

5.4 From search to set

Each row in the result table carries a single per-row action:

- **Add** — drop the concept into the concept set currently being edited. The row's button flips to **Remove** once it is in the set. **Add** is always enabled: when no concept set is open it auto-creates a new untitled set and opens the editor, confirming with a snackbar reading "Added '...' → New concept set".

The per-concept Include / Include descendants / Include mapped / Exclude flags described in older ATLAS releases are not set from the search results page; they are toggled per row inside the concept-set editor itself (Chapter 6).

Note

The exact set of vocabularies you can search is whatever your WebAPI was loaded with from *Athena* [10], the OHDSI vocabulary repository. Two analysts on the same ATLAS v3.0 release but with different vocabulary snapshots can see slightly different concept IDs; record the snapshot date along with your study.

5.5 Search strategies that work in practice

- Start broad, then narrow. Search for the disease by name, pick the highest-level standard concept that still represents what you mean, and rely on *include descendants* rather than listing every variant by hand.
- Cross-check against codes. If you know the ICD-10 chapter for your phenotype, search by code as well; the standard concept you arrive at should be the same.
- Use the source-specific record counts. A concept with zero records in your target source is unlikely to drive cohort entry on it, regardless of how well it represents the disease in the abstract.

Concept sets

A *concept set* is a named, reusable collection of concepts together with rules for how each concept participates. Concept sets are the building blocks of cohort definitions and most analyses; if you find yourself listing the same concepts in two cohorts, you should probably be referencing the same concept set in both.

6.1 Why a list of concepts is not enough

Two ideas separate a concept set from a flat list:

Descendants A standard concept usually has descendants in the OHDSI hierarchy. *Type 2 diabetes mellitus* (SNOMED) has many specific subtypes underneath it. With **Include descendants** ticked on the parent, every subtype is included automatically and stays current as the vocabulary evolves.

Mapped source codes Real patient records are often coded in source vocabularies (ICD-10, RxNorm, NDC). With **Include mapped** ticked, the set captures records that map up to the chosen standard concept even if they were originally recorded under a source code.

A concept set therefore expands at runtime into a — usually larger — list of concept IDs that the cohort SQL actually uses.

6.2 Creating a concept set

1. Open **Concept Sets** from the top navigation. The **Concept Sets** tab is the default; you land directly on the list.
2. Click **New Concept Set**. The editor opens as a side drawer on the right.
3. Give the set a descriptive name in the inline-edit title at the top of the drawer — research groups frequently adopt a convention such as [Phenotype] [author] [version].
4. A new set opens on the **Search** tab. Use the vocabulary search (Chapter 5) to find your starting concepts; click **Add** on each row.
5. Switch to the **Selected** tab to see the concepts in the set and adjust per-row flags.




Figure 6.1. Concept-set editor drawer with the **Included Concepts** / **Search** / **Recommend** / **Compare** tabs, the per-row Descendants / Mapped / Exclude flags, and the **Paste IDs** action top-right.

6.3 The editor tabs

The editor drawer has six tabs, in order: **Selected**, **Included**, **Source Codes**, **Search**, **Recommend**, and **Compare**. A new set opens on the **Search** tab so you can start adding concepts immediately; an existing set opens on **Selected**, where its concepts and per-row flags are shown.

6.4 The Selected tab

The **Selected** tab lists every concept you have explicitly added to the set, with three per-row toggles and a remove action:

Descendants Whether the set also covers concepts below this one in the OHDSI hierarchy.

Mapped Whether non-standard source codes that map up to this concept count as members.

Exclude When on, members covered by this row are subtracted from the set, even if another row would otherwise include them.

Membership itself is implicit: every row in the table is part of the set unless its **Exclude** flag is on. To take a concept out of the set entirely, use the row's remove action rather than an Include toggle.

These toggles compose. A common shape is "include the parent disease with descendants, then add an exclude row for one specific subtype". The **Selected** tab also has a

filter strip at the top of the table to narrow the displayed rows by domain, vocabulary, standard-concept flag, or exclusion status.

The concept-name cell on each row is a link to **Concept Details** (Chapter 5, Section 5.3); use it to inspect attributes, the hierarchy mini-map, and per-source prevalence without leaving the editor.

6.5 The Included tab

The **Selected** tab shows what you *added*; the **Included** tab shows what the set actually *resolves to*. It is a read-only table of the fully-expanded concept list — the parents you selected plus every descendant and mapped concept the toggles pull in, minus any excluded rows. There are no per-row toggles here; to change what is included, return to the **Selected** tab. When the set is empty the table shows an empty state reading "Add concepts on the Selected tab to see them here".

This is the list the cohort SQL ultimately uses, so it is the right place to sanity-check that a broad parent or a descendant subtree has not swept in more than you intended.

6.6 The Source Codes tab

The **Source Codes** tab lists the non-standard *source* codes that map to the set's included concepts — the ICD-10, Read, NDC or other source-vocabulary codes a record might carry before it is mapped up to a standard concept. The table has columns ID, Code, Name, Class, Domain, and Vocabulary. Reviewing it is the practical way to see which real-world codes the set will actually match in an ETL'd database.

6.7 Bulk-adding concepts

The right side of the editor's tab strip carries three import buttons — **Paste IDs**, **Import codes**, and **Import JSON** — for adding many concepts at once instead of one search result at a time.

6.7.1 Paste IDs

Paste IDs opens a dialog where you paste a list of concept IDs (one per line, or space- or comma-separated); ATLAS resolves them against the vocabulary and reports a matched / unmatched summary. This is the fastest way to import a code list from a spreadsheet or a phenotype paper.

6.7.2 Import codes

Import codes accepts *source* codes rather than concept IDs — for example ICD-10 E11.9. ATLAS resolves each code to the standard concept it maps to and adds those concepts to the set, again reporting what matched and what did not. Use it when your starting list is expressed in a source vocabulary rather than in OMOP concept IDs.

6.7.3 Import JSON

Import JSON accepts a concept-set expression of the form `{"items": [...]}` — the same shape ATLAS uses to serialise a set. Each item is added together with its per-row flags (descendants, mapped, exclude), so this round-trips a set exported from elsewhere while preserving how each concept participates.

A concept set that resolves to thousands of included concepts is rarely what you wanted. Either narrow the parent concept, turn off descendants, or add specific exclusions until the resolved list reads sensibly.

Before saving a concept set you intend to publish, sit with a domain-aware reviewer in front of the **Selected** tab and the record-count columns. This is the catch step for two common mistakes: (i) a standard concept whose mapped source codes are clinically broader than your phenotype intends (the up-hill mapping problem), and (ii) a descendant subtree that sweeps in obsolete or rarely-used codes a reviewer would have flagged. The complementary check — *orphan codes*, source codes a clinician would have included but the set does not capture — is what the HADES CohortDiagnostics package automates [28, 30].

6.8 Tags

The editor header carries a **Tags** button with a count badge. Clicking it opens a tag-selection dialog where you attach one or more tags to the set; the tags are saved together with the concept set. On the **Concept Sets** list they render as coloured chips next to each set's name, and they are one of the facets you can filter the list by (Section 6.13). Tags are the lightweight way to group related sets — by study, by domain, by review status — without baking that grouping into the set name.

6.9 Recommend

The **Recommend** tab uses PHOEBE 2.0 to suggest additional concepts that are commonly co-included with the ones you have already added. Recommendations need seed concepts to work from, so add at least one concept on the **Selected** tab before opening it; with an empty set there is nothing to recommend against. PHOEBE relies on co-occurrence patterns drawn from a large network of OMOP databases, so a concept it suggests is one that other phenotype authors have routinely paired with yours. Review each suggestion before accepting it — the recommender does not know your phenotype's intent, only the company it tends to keep in published code lists.

The results table is selectable: each row carries **Descendants**, **Mapped**, and **Exclude** checkboxes, and an **Add Selected (n)** button adds the ticked rows to the set with those




Figure 6.2. The tag-selection dialog, opened from the Tags button in the concept-set editor header, where tag groups are searched and applied to the set.

flags in one action. A **Refresh** action recomputes the recommendations after you have changed the set's concepts.

The tab is unavailable when the WebAPI's vocabulary schema does not include the PHOEBE 2.0 tables; the panel itself surfaces this condition.

6.10 Compare

The **Compare** tab compares the concept set you are editing (**Concept Set 1**) against another saved concept set (**Concept Set 2**). The workflow is:

1. Click **Choose** on the toolbar; pick the second set from the picker dialog.
2. Click **Compare Concept Sets**. ATLAS resolves both sets, computes the intersection and the two differences, and renders a Venn diagram alongside a result table.
3. A segmented toggle above the results chooses what is compared: **Expression concepts** compares the concepts as explicitly listed in each set's expression, while **Included concepts** compares the fully-resolved expansions (descendants and mapped codes included). The two can differ markedly when either set leans on **Include descendants**.
4. The result table has a **Match** column that tags every concept as *Both* when it appears in both sets, or with the *name of the set* that uniquely contains it, alongside the usual concept ID, code, name, domain, vocabulary and class columns. Sort or scroll the set-specific rows to read off exactly which concepts moved when you compare two phenotype iterations.

5. **Export** on the toolbar writes the result table to a CSV file for review or attachment to a phenotype submission.

Use the comparison to:

- review what changed between two phenotype iterations before retiring the older one — the rows tagged with the older set's name are what it captured that the new one no longer does;
- check whether two phenotypes that should be disjoint (e.g. a positive and a negative control) actually are — the Venn diagram should have an empty intersection;
- migrate from a peer-curated library set to a local one without losing terms accidentally.

A worked iteration: you have set v1 published in a study, set v2 in development. Open v2, switch to **Compare**, choose v1 as **Concept Set 2**. The rows tagged with v1's name are the terms v2 dropped — defend each one in the change log. The rows tagged with v2's name are the new terms in v2 — justify each. The *Both* rows are the unchanged backbone. Do this review before you point any cohort at v2.

6.11 Saving

Save (or **Create** on a new set) writes the current concept set to WebAPI. The first save creates a permanent ID; later saves create new versions (Chapter 14). **Delete** removes the set; both buttons are disabled when you do not hold the relevant per-entity permission. The editor header also carries a **Versions** button, which opens the version history for the set (Chapter 14).

Include Mapped pulls in records whose *source* code maps to your standard concept, even when the mapping is "up-hill" to a broader parent than the source code itself denotes. For precision-sensitive phenotypes, leave it off and rely on standard concepts the source database was ETL'd to use directly. The mapping mechanics are covered in *The Book of OHDSI*, chapter 5 [20].

6.12 Starting from a phenotype library

Before authoring a concept set from scratch, check whether the phenotype already exists in a community library. Several peer-curated libraries can be searched first:

- *The OHDSI Phenotype Library* [17] holds peer-reviewable cohort definitions in OMOP-CDM ATLAS JSON. Each entry has a stable identifier and each release of the library carries its own DOI for citation.

- *The HDR UK Phenotype Library* [3] indexes more than a thousand phenotype algorithms expressed as clinical code lists — ICD, SNOMED, Read — across cardiology, oncology, COVID-19 and other areas, contributed by groups across UK research; it is API-accessible and ships an R client.
- *EHDEN* [2] (the European Health Data & Evidence Network) curates phenotypes used in pharmacoepidemiology studies on its federated 100M-patient OMOP estate, and produces training material on phenotype definition, characterisation, and evaluation that complements the OHDSI tools.

A concept set or cohort imported from any of these is almost always cheaper to adapt than to author — and visibly more defensible at peer review — than starting with a blank expression.

When you reuse a definition from the OHDSI Phenotype Library, record both the cohort identifier and the library release version (each release has its own DOI). Cohort definitions accepted into a release are immutable for that release, but a definition can be marked deprecated [D] or in error [E] in a later one. Pinning the release insulates your study from those later changes.

6.13 Discovery

Concept sets accumulate quickly. The list page on the **Concept Sets** tab carries a free-text search at the top that matches against the set name, and next to it a **Filters** menu offering **Author**, **Tags**, and **Created** / **Modified** date-range facets. Each active facet shows below the bar as a removable chip, so you can stack, say, an author and a tag and see at a glance what is narrowing the list. The table columns are ID, Name, Tags (rendered as coloured chips), Created, Updated, and Author, and they are sortable. Adopt a naming convention — e.g. [Phenotype] [author] [version] — and tag sets consistently so both the search and the facets stay useful as the library grows.

PART III

Cohort Definitions

Cohort definitions

A *cohort* is a set of persons who satisfy one or more criteria for a duration of time. A *cohort definition* is the formal, machine-executable description of those criteria. This chapter is the longest in the manual because the cohort builder is where you spend most of your time in ATLAS v3.0.

The cohort builder is divided into numbered sections matching the logical pieces of a definition — entry event, inclusion criteria, exit, and censoring — with a collapsible Generation section at the bottom for running the cohort against connected sources.

Note

A cohort definition built in ATLAS is one stage of a longer phenotype-development workflow that includes diagnostics, evaluation, and submission to a community library. The methodology is documented in *The Book of OHDSI* [20] (chapters 10, 16, 18) and is out of scope for this manual; we focus on the ATLAS controls.

ATLAS's cohort model lets one person belong to the same cohort during *multiple non-overlapping time periods* — a person who initiated metformin in 2014 and again in 2019 has two cohort entries. They are never in the cohort twice simultaneously. This matters when you read reports: the "cohort entries" count is greater than or equal to the "persons" count, and the difference tells you how often re-entry happens in your data.

If you came from SAS or R, you are looking for NOT EXISTS. ATLAS does not have it. Every "absence" constraint — "no prior exposure to drug X", "no diagnosis of Y" — is expressed as an inclusion rule with cardinality *exactly 0 occurrences* in the relevant temporal window. This pattern recurs throughout the chapter; the cohort builder itself has no negation operator.

7.1 Live preview

When the deployment has TrexSQL enabled and a per-source patient cache built (Chapter 15, "Caches and TrexSQL"), the cohort builder shows running patient counts as you edit — no **Generate** round-trip required. There is no longer a single "patient count" bar

Figure 7.1. The cohort builder. A numbered step rail walks down the page (Cohort Entry Events, Inclusion Criteria, Cohort Exit & Eras, Generation); the toolbar carries Cancel, Export, and Save.

at the top of the builder; counts appear *inline*, one per inclusion rule, in the inclusion-rule rail described in Section 7.5.

The data source the live preview runs against is chosen from the **Live preview cache** selector at the top of the inclusion-rule rail (Section 7.5), sitting directly above the **Entry events** count row. The selector lists every connected source with a status chip:

Ready. The cache is up to date; live counts are current.

Stale. The cache exists but is older than the underlying data. Live counts still display but with a stale-tone chip; rebuild the cache for accuracy.

Building. The cache is being built. Live counts appear once the build finishes.

Error. The last cache build failed. Live counts are unavailable for this source until it is rebuilt.

Not built. The source has no cache. Ask an administrator to build it from the configuration panel's **Cache** tab; live counts are unavailable until then.

A source picked here is also where the inclusion rail's **Entry events** count, the per-rule funnel counts, and the **Qualifying cohort** tile come from.

A rule that drops the live count from 80% to 0% the moment you tick it is almost always a misconfigured concept set or an inverted cardinality, not a real epidemiological signal. Treat the inline counts as your in-place sanity check: they catch mistakes that would otherwise only surface after a full **Generate**.

Note

The selector and the inline counts are hidden entirely on deployments where TrexSQL is not enabled. If you do not see them, the deployment runs without live preview — proceed with **Generate** as the only feedback loop.

7.2 Listing and creating cohorts

Open **Cohort Definitions** from the navigation. The list page supports two views — tile and table — toggled in the top-right. Cards show the cohort name, author, last modification date, tags, and quick actions. Use the search bar to filter by name; sort by recency to find what you were last working on.

- **New Cohort** starts an empty definition.
- **Import** accepts a cohort definition JSON file (the format ATLAS writes when you export). ATLAS v3.0 will run it through its ATLAS-2.x converter where needed before the editor opens.
- Clicking a card opens that cohort in the editor.

7.2.1 Duplicating a cohort

Every cohort card (and every row in the table view) carries a **Duplicate** action. It copies the full definition under the name "*<name> (copy)*" — and, when that name already exists, auto-increments to "*(copy 2)*", "*(copy 3)*", and so on — then opens the new copy directly in the builder. WebAPI has no server-side copy endpoint for cohorts, so ATLAS builds the duplicate client-side from the fetched definition; the copy is an independent cohort with its own ID and version history.

7.3 The four logical pieces of a definition

Every cohort definition answers four questions, in this order:

1. **Who enters the cohort, and on what date?** — entry event (also called primary criteria, or initial event).
2. **Of those entrants, who do we keep?** — inclusion criteria with cardinality and temporal logic.
3. **When does the cohort end for each person?** — exit criteria.
4. **What ends membership early?** — censoring events.

The rest of this chapter walks each in turn.

7.4 Entry event (primary criteria)

The entry event defines what kind of clinical record makes a person a candidate for the cohort, and which date of that record becomes the person's *index date*. You add at least one entry event from one of these criterion types:

- Condition occurrence, Condition era
- Drug exposure, Drug era, Dose era
- Procedure occurrence
- Measurement
- Observation, Observation period
- Visit occurrence, Visit detail
- Device exposure
- Death
- Specimen
- Payer plan period
- Demographic, Location region

When you add an entry event, the editor presents a panel with a concept set picker followed by a list of *attribute filters* appropriate to the domain. For a condition occurrence, those attributes include age at occurrence, gender, condition type, condition status, and visit type. For a drug exposure, they include dose, quantity, days supply, refills, route, and stop reason.

Figure 7.2. Entry-event panel. The **Cohort entry on** toggle on the right (earliest event / all / latest event) is the per-section qualifying-events limit; the observation-period chip ("365 days Before and 0 days After the event index date") is the look-back / follow-up requirement.

7.4.1 Cohort entry observation window

Below the entry event you set how many days of continuous observation you require before and after the index date.

- *Days before index* controls washout. A value of 365 means "the person must have been observable in the data for at least a year before they entered the cohort", which is the usual way to require a clean look-back period.
- *Days after index* ensures a minimum follow-up.

Encode the required look-back as *Days before index* on the entry event itself, not as an inclusion rule. Doing it as an inclusion rule biases downstream incidence and characterization denominators; the OHDSI Themis convention is the entry-event form [20, 25].

7.4.2 Three qualifying-event limits

ATLAS exposes three independent limits that decide *which* qualifying events count, each as an **earliest event** / **all** / **latest event** toggle:

Cohort entry on applies to the entry-event panel itself. **earliest event** produces a *new-user* or *incident* cohort (only the person's first qualifying record); **all** lets every qualifying record create a cohort entry; **latest event** keeps only the most recent.

with applies to the optional **Restrict initial events** block that further restricts entry events *before* the inclusion rules below run. Same three values; controls which qualifying entry the additional criteria are evaluated against.

Apply rules to sits above the inclusion-rule list and decides whether the rules below apply to the first, every, or last qualifying record per person. With **earliest event**, ATLAS narrows to the earliest qualifying entry and applies the inclusion rules to that one event — usually what new-user designs want.

All plus generous descendants on a chronic-disease concept set is a fast way to multiply rows by an order of magnitude. Pick the limit on each section that matches the question.

7.5 Inclusion criteria

Inclusion criteria are applied after entry events. Each rule is itself a small cohort criterion — "at least one diagnosis of X in the 365 days before index", "no exposure to Y any time before index" — and the person must satisfy every rule that you mark as required.

Note

Each inclusion rule you add raises specificity at the cost of sensitivity — the trade-off is the same as in any phenotype algorithm. Inspect the inclusion-rule report and a sample of profiles (Chapter 12) to see which side of that trade you have landed on. The methodology is covered in *The Book of OHDSI* [20], chapter 10.

7.5.1 Rail and detail layout

The Inclusion Criteria section is a master–detail layout. The left pane is an *inclusion-rule rail* listing every rule; the right pane is the editor for the selected rule. Click a rule in

the rail to load it on the right; **Add rule** at the top of the rail appends a new one and selects it.

When live preview is on (the deployment has TrexSQL enabled and the source picked in **Live preview cache** has a built cache, Section 7.1), the rail carries running counts:

- An **Entry events** row at the top reads " N of T dataset patients $\cdot p\%$ " — how many of the source's persons satisfy the entry event before any inclusion rule applies.
- Each rule row shows its surviving count, the percentage of entry events retained at that step, and a horizontal fill bar sized to that percentage. Rule rows turn amber as retention drops; the rule currently being recomputed shows a small spinner (recomputation is debounced 500 ms after the last edit).
- A **Qualifying cohort** tile at the bottom of the rail reads off the final count and the percentage of entry events that survive every rule.

Rules in the rail can be reordered by drag-and-drop using the handle at the left of each row; a drop indicator highlights the target position. Order matters for the attrition story even though it does not change the final cohort: the per-rule count above is the count after that rule and every previous rule has been applied.

The right pane carries a sticky header with the selected rule's name plus two icon buttons — **Edit name** (rename the rule) and **Delete** (remove it). The rule's body sits below the header.

7.5.2 Anatomy of a rule

1. **Domain.** Pick what kind of event the rule is about (condition, drug, procedure, measurement, observation, visit, device, era, etc.).
2. **Concept set.** Reference an existing concept set or open a popup to build one inline.
3. **Attribute filters.** Same domain-specific list as for the entry event.
4. **Cardinality.** Choose *at least*, *at most*, *exactly*, plus a count N . A third control, **Counting method**, decides what N actually counts: *all occurrences*, *distinct concepts*, *distinct start dates*, or *distinct visits* (this maps to CIRCE's `Occurrence.CountMethod`).
5. **Temporal window.** Specify the time relative to index in which the events count, e.g. *Days: 365 days before to 0 days before*, with *Use index date* or *Use index + cohort exit* as anchors.

7.5.3 Grouping criteria with AND / OR

Each inclusion rule can hold more than one criterion. The **Add criteria** button under a rule adds another criterion to the same group. The combinator that decides how the criteria combine is *not* a dropdown at the top of the rule: it is a clickable, colour-coded

vertical label running down the left edge of the group card. Clicking it opens a small segmented popover with four options:

All (logical AND) — the person must satisfy every criterion in the group.

Any (logical OR) — the person must satisfy at least one.

At least N of — the person must satisfy any N of the criteria, where N is configurable.

At most N of — the person must satisfy no more than N of the criteria.

The **Add criteria** menu also offers a **Group** entry that inserts a nested *sub-group* with its own combinator, which is how complex phenotypes like "diabetes diagnosis AND (insulin use OR HbA1c ≥ 6.5)" are expressed. Build the nesting carefully — two criteria at the same level combine differently from one criterion plus a sub-group.




Figure 7.3. Expanded inclusion rule. Inside the rule, criteria groups carry a cardinality ("at least 1") and each criterion carries a concept set, attribute filters, and a temporal-window chip ("30 days before to 0 days after index start").

A phenotype that requires a confirming second event after index introduces *immortal time* between the first and second records. Either index on the second event in ATLAS (so the immortal window precedes index) or keep the confirming window short. The bias is documented in Suissa [36] and Swerdel et al. [39].

7.5.4 The criteria event card

Every individual criterion — in the entry event, in the **Restrict initial events** block, and inside an inclusion rule — is rendered by the *same* shared card component, so the controls look and behave identically wherever you meet them. The card has two consistent features:

- A *cardinality sidebar* down the left edge. Click it to set **At least**, **Exactly**, or **At most** plus a count for how many matching events the person must have.
- A *temporal-window chip* that states the time window, relative to index, in which those events count.

The one difference is context-driven: on the entry event the cardinality sidebar is hidden (an entry event is the anchor, not a counted criterion), while inclusion and nested criteria show it. Learning the card once means you have learned it for all three sections.




Figure 7.4. The shared criteria event card — here an inclusion-rule criterion — showing the left-edge cardinality sidebar (**at least/exactly/at most** plus a count) and the **Add Temporal Window** control.

7.5.5 Reading the inclusion report

After generation, the **Inclusion report** action on a row of the generation table (Section 7.8) opens the inclusion-rule report in a side drawer. The report has four parts, all keyed to the same data:

Mode tabs. **By Person** (default) counts each person once across all matching cohort entries; **By Event** counts every qualifying entry separately. The two views agree exactly when the cohort has **Cohort entry on** set to **earliest event**; they diverge as soon as a person can have multiple entries.

Summary stats. A row of stat cards: **Match rate**, **Matches** (the final cohort size), **Lost** (entry events removed by the rules), and **Total events**.

Attrition. The **Attrition** panel is a per-rule horizontal funnel chart with one bar per inclusion rule, sized to the percentage of entry events surviving that rule. Bars colour-grade from green through amber to red as retention drops; each bar carries the absolute count and percentage, with large counts shown in compact SI form (e.g. 1.2M). A footer summarises the journey: *"Initial: N patients → Final: M patients (X.X% retained)"* with the retention figure tinted by severity. The panel header offers a **Download CSV** action (the file is written as `attrition.csv`) and a **Save as PNG** export of the chart. When the cohort has no inclusion rules or has not been generated yet, the panel shows an empty state instead.

Per-rule satisfaction. A table with one row per rule, showing how many entry events each rule kept or dropped. The **Cumulative remaining** column is the funnel's numbers in tabular form — how many entry events survive after this rule and every previous rule have been applied.

Population breakdown. A treemap of the rule-satisfaction signature: each tile is a unique combination of which rules a person passed and which they failed, sized to the number of people with that signature. Hover for the bitfield. The treemap surfaces the dominant failure pattern at a glance — if one tile dominates, that combination is the one to investigate.

When counts surprise you, read the funnel as a sequence of suspects:

- A rule that drops to near-zero on its own usually means the concept set resolved to nothing on the source — check the concept set's **Included Concepts** tab and re-run on a source with known coverage.
- A rule that drops nobody is either trivially satisfied (the temporal window is too wide, the cardinality is too lax) or redundant with an earlier rule. Either way it is not pulling weight.
- A rule that drops disproportionately many people compared with a peer source signals an ETL or vocabulary-mapping difference — the same rule against the same JSON should drop comparable fractions. The HADES CohortDiagnostics package formalises this cross-source comparison with attrition tables side by side [30].

7.6 Cohort exit and censoring

7.6.1 Exit criteria

Exit defines the per-person end date. Choose one of:

End of continuous observation. The person leaves the cohort when they leave observation in the source. The most common choice for prevalent or chronic-disease cohorts.

Fixed duration. A specified number of days after index. Useful for time-to-event analyses with a hard horizon.

End of continuous drug exposure. Used for drug-utilisation cohorts. You configure the persistence window (the gap between consecutive exposures still considered the same era) and the surveillance window.

Event criteria. Exit on the first occurrence of a specified event after index.

7.6.2 Censoring events

Censoring events end the person's membership early when an outcome makes continued follow-up nonsensical. Common censoring choices include death, end of insurance enrolment, and onset of an outcome you want to exclude.

7.6.3 Era collapse

For cohorts where the same person can enter and re-enter, you can **Collapse cohort entries into eras** with a configurable gap days. With a 30-day gap, two cohort entries

20 days apart are merged into one continuous era; with a zero gap, every entry stays separate.

The 30-day persistence window is the OHDSI convention for drug eras because it has historically performed well for drug–outcome estimation, not because it is biologically motivated. For long-acting injectables and depot formulations, 90 days or more is common; for acute treatments where a real treatment break is clinically meaningful within days, shorter values matter. Treat the window as a study parameter: report it in the manuscript and, on sensitive analyses, run the cohort with ± 30 days around your choice as a sensitivity check.

Figure 7.5. Cohort Exit & Eras. The strategy toggle picks Observation, Fixed duration, or Drug exposure; Censoring Events end membership early; the era collapse gap (in days) merges adjacent cohort entries.

7.7 Concept sets used by this cohort

The **Concept Sets** tab inside a cohort definition lists every concept set the definition references. From here you can review the resolved included concepts without leaving the editor. Renaming a concept set here renames its inline copy on this cohort but does not affect the shared library version — that is handled through the version history (Chapter 14).

Note

A cohort definition stores its own *copy* of every concept set it references. This copy travels with the cohort JSON when you export it, so a downstream user can reproduce the cohort even without access to your concept-set library. The cost is that updating a concept set in the library does *not* retroactively update the embedded copy on every cohort that uses it — re-saving the cohort against the library version is a deliberate, explicit action.

7.8 Generation

Generation runs the definition's SQL against a chosen source. The cohort builder carries a collapsible **Generation** section below the cohort body — there is no longer a toolbar **Generate** button or a separate side panel.

The section header summarises status at a glance:

- a status chip on the right reading **Save cohort to generate, No generations yet**, or *X / Y generated* (green when all sources are done, amber while a run is pending or in progress, red when any run has failed);
- a **Last run** · *⟨relative time⟩ on ⟨source⟩* meta line for the most recent successful run;
- an **Only generated** toggle that filters the source table down to sources the cohort has already been generated against, hiding the rest;
- a **Generate all** primary button on the right that runs the cohort against every source you have write access on, skipping sources whose runs are already pending or running.

The body is a single table with one row per connected source and the following columns: **Data source, Status, Patients, Last run, Duration, Actions**. Per-row actions:

Run starts a new generation for that source; while a run is pending or in progress the row's status chip animates and the action is disabled.

History icon opens the **Previous runs** dialog, listing every run for that source with timestamps and durations, so you can compare successive generations.

Inclusion report opens the inclusion-rule report in a side drawer (Section 7.5.5). Disabled until the row's last run has completed.

Samples opens the cohort samples panel for that source in the same drawer (Section 7.8.1). Also disabled until the run has completed.

Both **Inclusion report** and **Samples** reflect their state in the URL (*?report=inclusion&source=⟨key⟩* or *?report=samples&source=⟨key⟩*), so a link to a particular cohort with a report open is shareable.

Note

A row whose source you do not have write access on shows a disabled **Run** button with the message *"You don't have write access on this source."* The row's read-only data — last-run status, patient count, history — remains visible.

7.8.1 Cohort samples

The **Samples** panel is the fastest path from a generated cohort to individual patient review. Click **New sample** at the top of the panel to open the sample-creation dialog; specify:

Sample name a descriptive label, e.g. `review-2025-01-women-over-65`.

Number of persons how many to draw, up to the per-source maximum.

Gender (optional) one or more of **Male**, **Female**, **Other** / **non-binary**.

Age (optional) a comparator (*less than*, *greater than*, *between*, etc.) plus a value or range.

ATLAS runs a server-side random draw constrained by the filters and lists the resulting persons. Selecting a sample loads its detail block underneath; clicking a person from there opens their **Profile** *in the same drawer*, with a back-arrow in the drawer header that returns to the samples list. **Refresh** re-rolls the same sample with a new draw, **Delete** removes it.

Use a sample of 20–50 patients as the lightweight half of phenotype validation: skim each profile, note any obvious mismatches, and adjust the cohort definition. This is much cheaper than a full PheValuator run (Chapter 12, "Going further") and catches the gross errors first.

Note

Generation is the only step that actually queries patient data, and on large sources it can take a while. It is also the first place a flaw in the definition — a concept set that returns nothing, a temporal window that excludes everyone — becomes visible. The inclusion-rule report (Section 7.5.5) is your first stop when counts surprise you.

7.9 Going further: external diagnostic tools

The reports built into the Generation section are deliberately a quick view — attrition by rule, plus a sample. For the deeper diagnostics expected of a sharable phenotype, two external HADES R packages consume the JSON ATLAS exports:

CohortDiagnostics [28, 30] produces cross-source attrition, index-event breakdown, source-code review, orphan codes, incidence by age / sex / year, cohort overlap, and time distributions. Fast, runs without a gold standard.

PheValuator (Chapter 12, Section 12.5) estimates sensitivity, specificity, and PPV against a probabilistic gold standard.

Both are recommended evidence for OHDSI Phenotype Library submissions [26]. The methodology behind them lives in *The Book of OHDSI* [20], chapters 16 and 18.

For *negative-control outcomes* — the supplementary cohorts many OHDSI designs use to detect residual bias — ATLAS is just where you build them. Tag them `negative-control`; the methodology and recommended counts are covered in the Book of OHDSI's Method Validity chapter.

7.10 Saving, copying and exporting

Save the cohort with a meaningful name as soon as you start editing, even before any criteria are filled in. The first **Save** creates a permanent ID and unlocks version history; subsequent edits then accumulate as versions you can roll back to, rather than as one giant undo stack you can lose.

- **Save** writes the cohort to WebAPI. Repeated saves create new versions (Chapter 14). The button is disabled when you do not hold write permission on the cohort.
- The **Export** action on the toolbar opens a small menu with two options: **Download as file** writes the definition's JSON to disk, and **Copy to clipboard** copies the same JSON for pasting elsewhere. The same JSON can be re-imported into this ATLAS instance or another.

7.10.1 Tagging a cohort

The cohort toolbar carries a tag icon (MDI-TAG-MULTIPLE) with a count badge of the current tags. Click it to open the **Manage Tags** dialog, which lists every tag defined in the deployment grouped by tag group; tick the ones that apply.

Tags are deployment-wide artefacts — they are created and curated from the configuration panel's **Tags** tab (Chapter 15), not from the cohort. A useful convention is one tag per project area (e.g. cardiology, oncology), one per study phase (exploratory, published, negative-control), and a small set of reusable qualifiers (phenotype, drug-class, outcome). Tags appear on cohort cards in the list page and are searchable by sorting on the tag column.

The same tag mechanism applies to concept sets (Chapter 6), characterizations (Chapter 9), pathways (Chapter 10), and incidence rates (Chapter 11); the icon, dialog, and behaviour are identical across editors.

7.10.2 Auto-save and draft recovery

ATLAS auto-saves a draft of the cohort to the browser's session storage every 30 seconds while the editor is dirty, even before the first deliberate **Save**. The mechanism matters in three situations:

- **Token expiry mid-edit.** If your session expires and you log back in, the draft is restored on return to the editor.
- **Accidental tab close.** Closing the tab and reopening it within the same browser session restores the draft.
- **New cohort that was never saved.** The draft is keyed to the editor, not to a cohort ID, so an unsaved new cohort survives a refresh.

The draft is *not* synced to other browsers, devices, or tabs (it is per-tab), and it is cleared when the cohort is deliberately saved or the browser session ends. Treat it as a crash-recovery net, not a substitute for **Save**.

PART IV

Analyses

The Analysis Hub

The four analyses in this part of the manual — characterizations, cohort pathways, incidence rates, and feature analyses — no longer live behind four separate navigation entries. ATLAS v3.0 gathers them into a single *Analysis Hub*: one page with a tab per analysis type. This chapter explains how to reach each analysis and how the hub is laid out, so the chapters that follow can concentrate on the analyses themselves.

8.1 Where the analyses live

The top navigation bar carries four destinations: `Data Sources`, `Concepts`, `Cohorts`, and `Analysis`. Opening `Analysis` lands you in the hub, which presents its four analyses as tabs across the top of the workspace:

Feature Analyses. Reusable covariate definitions — the building blocks that characterizations draw on (Chapter 13).

Characterizations. Descriptive summaries of one or more cohorts (Chapter 9).

Pathways. Sequential treatment- and event-pattern diagrams (Chapter 10).

Incidence Rates. Outcome rates within a target population over a time-at-risk window (Chapter 11).

The hub remembers the tab you used last and reopens it the next time you return; the first time you arrive it defaults to **Feature Analyses**. Each tab shows its own list of saved analyses, and each list carries its own **New** button.

Note

Old bookmarks still work. The former standalone URLs — `/characterizations`, `/feature-analyses`, `/pathways`, and `/incidence-rates` — redirect into the matching hub tab, so links you saved against an earlier build land in the right place.

8.2 Lists, editors, and results

The hub is the home for the *lists* of saved analyses. Building or reviewing an individual analysis takes you out of the tabbed hub and into a full-screen editor or result page — creating a new characterization, opening a pathway's sunburst, or drilling into an incidence-rate workbench each fills the whole window rather than sitting inside a tab.

Closing or navigating back from one of these full-screen views returns you to the hub tab you started from.

This split keeps the hub itself uncluttered: it is a launcher and an index, while the detailed work happens on a dedicated page with room for the design rail, canvas, and insight panels described in the chapters ahead.




Figure 8.1. The Analysis Hub with its four tabs — Feature Analyses, Characterizations, Pathways, Incidence Rates — showing the saved-analysis list for the active tab and its New button.

8.3 A note on Profiles

Person *profiles* (Chapter 12) are an analysis in spirit, but they are deliberately *not* part of the hub and do not appear in the top navigation. A profile is always about one specific person, so you reach it in context — by following a **View profile** action from elsewhere in the application, or by opening a profile URL directly — rather than by browsing a list. Chapter 12 covers those entry points.

Characterizations

A *characterization* summarises the clinical features of one or more cohorts. Where a cohort report (Chapter 7) describes one cohort on one source, a characterization is broader: it can put several target cohorts side by side, attach outcomes, and apply a deliberate, reusable set of feature analyses.

9.1 What you can answer with a characterization

Typical questions:

- Among new users of drug A versus drug B, what does the patient mix look like at index — age, comorbidities, prior medications?
- In a population with disease X, how does the burden of related conditions evolve in the year after diagnosis?
- Across five different data sources, how comparable are the baseline characteristics of the same target cohort?

A well-designed characterization is what you publish alongside a study to demonstrate that “the patients in my cohort look like what you would expect of patients with this condition”.

Note

For comparative-effectiveness work, you typically run a characterization with two target cohorts (the *target* and the *comparator* of an active-comparator new-user design [42]) before propensity adjustment. The study design and methodology are covered in *The Book of OHDSI* [20], chapters 7 and 12.

Note

The ATLAS characterization presets give four standard time-window durations relative to the cohort entry (*index*) date: *Any time prior*, *Long-term* (365 days), *Medium-term* (180 days), and *Short-term* (30 days). A typical baseline bundle covers all four so reviewers can see how prevalence changes as the look-back window narrows. The underlying distinction between baseline, post-index, and time-at-risk windows is methodology covered in *The Book of OHDSI* [20], chapter 11.

9.2 Anatomy of a characterization

Open **Analysis** and switch to the **Characterizations** tab of the Analysis Hub (Chapter 8), then click **New Characterization**. The editor has three sections, in order:

Cohorts. One or more cohorts whose members you want to describe. There is no separate outcome-cohort slot — outcome-stratified characterisations are expressed by adding the outcome cohort here and a matching subgroup below.

Feature analyses. The list of covariates to compute. Each feature analysis is itself a reusable artefact (Chapter 13). A typical default set includes demographics, conditions in the year before index, drug exposures in the year before index, prior visits, and Charlson comorbidity components.

Subgroup analyses (optional). Overlapping selections of the population built like miniature inclusion rules: concept set, attribute filters, cardinality, and a temporal window relative to index. Each subgroup becomes an extra column in the result alongside the main cohort. Use them to ask questions like "what does the cohort look like *among those with prior MI?*" without having to redefine and regenerate the cohort itself. A **Calculate subgroup analyses only** toggle is available when you want subgroup columns without the full characterisation.

The data sources to run on are picked at generation time, not on the design form.

9.2.1 Building a subgroup

New subgroup opens a small editor that mirrors the inclusion-rule panel from the cohort builder: name the subgroup, reference or build a concept set, set attribute filters and a cardinality, and specify a temporal window relative to index. The subgroup membership is computed at characterization time and yields a column whose values are restricted to persons satisfying the criteria.

9.3 Generating and viewing results

The canvas shows a *data-source run table* at the top with one row per connected source. Each row carries the source name, its current execution status (**COMPLETED**, **RUNNING**, **FAILED**, **CANCELED**, or pending / starting / stopping while in flight), the last-run timestamp, the duration, and a single primary action button on the right whose label cycles through **Generate** (never run), **Cancel** (while running) and **Rerun** (after completion). A history icon (MDI-HISTORY) sits next to it with a small count badge when more than one run exists. WebAPI executes one characterization run per (cohort × source) pair; for each cohort it computes every feature analysis.

Click the history icon on a source's row to open the **Previous runs** dialog scoped to that source — it lists every historical run with its date, status, duration, and an eye icon to reopen completed results. The same dialog opens from the matching icon on the IR and pathway editors, so the navigation pattern is identical across the three analysis types.

When a run completes you see two output styles:




Figure 9.1. Characterization editor. The design rail on the left lists Linked Cohorts, Linked Feature Analyses, and Subgroup analyses; the canvas on the right shows the data-source run table above the result panels.

Prevalence-style features (binary covariates such as "had diabetes in the prior year"): a percentage with a count column, sorted by prevalence within each target cohort.

Distribution-style features (continuous covariates such as age): mean, standard deviation, median, and quartiles.

9.3.1 Result view modes

A toggle on the canvas toolbar switches the result between three view modes:

Per-analysis the default view — the prevalence and distribution tables described above, one block per feature analysis.

Baseline a computed "Table 1" summary that lays the cohorts out as columns and the covariates as rows, in the style expected of a published baseline-characteristics table.

Overview a charts dashboard rather than tables: box plots for the age distributions, demographic bar charts, a standardized-mean-difference (SMD) chart for spotting covariates that separate the cohorts, and a top-covariates bar chart of the most prevalent features.

A filter strip above the result table carries two multi-selects — **Domain** and **Analysis** — and a single-select **Cohort** dropdown that you combine to narrow a busy result to the rows relevant to the question at hand. Alongside them a **Threshold** $\geq$ prevalence filter hides every covariate whose prevalence falls below the value you set, so you can strip a long report down to its common features in one step.




Figure 9.2. The characterization result in Overview mode — a Distributions row of box plots (Charlson index, age) and a Demographic Profile of grouped bar charts — reached from the Per-analysis / Baseline / Overview toggle on the canvas toolbar.

Each row has an **Explore** action that opens a single-covariate detail dialog. The dialog leads with the covariate’s metadata — its analysis, domain, and concept id, the concept id linking out to the concept page — and then shows a per-cohort comparison table with an inline bar indicator against each cohort’s value, so you can read one covariate’s prevalence across all the target cohorts side by side. It is a focused look at a single row, not a drill-down into the concepts that contribute to it.

9.4 Export

Export CSV on the canvas toolbar dumps the entire result table for use in a spreadsheet, a manuscript table, or further analysis in R. The downloaded file is named `characterization-{executionId}.csv` (or `characterization-results.csv` when no specific run is selected) so you can match it back to the run that produced it.

9.5 Common pitfalls

- **Empty rows on a source.** If the target cohort generates zero persons there, every covariate is necessarily blank. Check the cohort generation log first.
- **Picking too many features.** Each feature analysis costs time and produces a row in the report. A scoped, reproducible set is more useful than an unfiltered dump.
- **Window mismatch.** Feature analyses define their own windows relative to index. If your cohort exit is shorter than a covariate’s window, the covariate may pull data from outside the person’s cohort membership.




Figure 9.3. The design rail in detail: a Linked Cohorts section, a Linked Feature Analyses section listing each linked feature with its type and stat type, and a Subgroup analyses section with the **New subgroup** action.

9.6 Reproducing a characterization in R

Once a characterization is saved, the same definition can be re-executed in R against a CDM the WebAPI may not be connected to. The HADES FeatureExtraction [\[15\]](#) package consumes the same JSON the ATLAS v3.0 editor produces; build a `covariateSettings` object via `createCovariateSettings(...)` or, for the aggregated form ATLAS computes, `createDefaultCovariateSettings(...)` with `aggregated = TRUE`. The R workflow is the right choice when the result needs to feed a downstream HADES analysis or when you want to add custom covariates not expressible in the ATLAS UI.

Cohort pathways

Cohort pathways answer the question: in what order do clinical events happen for the people in my cohort? You feed in a target cohort plus a small set of *event cohorts* (each representing a treatment or a clinical milestone), and ATLAS v3.0 produces a sunburst-style diagram that visualises the most common sequences. The technique was used at network scale to characterize treatment patterns across the OHDSI collaborators in Hripcsak et al. [5].

10.1 When this analysis is the right tool

Pathways shine when the comparison of interest is sequential rather than cross-sectional. Examples:

- Among newly diagnosed type 2 diabetes patients, what is the usual order in which oral antidiabetic drug classes are tried?
- For patients with depression, how often is the second prescribed medication a different class from the first?
- In oncology, what does the line of therapy distribution look like?

10.2 Setting up a pathway

Open Analysis and switch to the **Pathways** tab of the Analysis Hub (Chapter 8), click **New Pathway**, and fill in:

Target cohorts. One or more cohorts whose pathways you want to compare. The result page lays out one sunburst per target.

Event cohorts. Each event cohort represents one node type in the pathway — typically one drug class or one milestone per cohort. The shorter and more orthogonal this list, the more readable the resulting diagram.

Collapse window (days). Two event cohorts that overlap inside this window are treated as a combination at the same step rather than as a sequence. Pickable values: 1, 3, 5, 7, 10, 14, 30.

Minimum cell count. Sequences with fewer than this many persons are suppressed (selectable from 1 to 10). Pair this with whatever small-cell suppression policy your project or IRB requires so the diagram never reveals near-unique patient trajectories.

Maximum path length. How many steps the analysis follows past the index event, from 1 to 7. Longer paths yield more detail but produce sunbursts that fragment into increasingly small wedges and rapidly cross the minimum cell count threshold.

Allow repeats. When on, the same event cohort may appear more than once on a single person's pathway (e.g. a drug re-initiation after a gap). When off, the second occurrence is collapsed into the first.

Targets are typically *treatment-initiator* cohorts (e.g. "new users of any first-line anti-hypertensive"), and event cohorts are *one per intervention* (one per drug class, one per procedure type) [5]. Keeping the event cohorts orthogonal and small in number is the single biggest determinant of whether the resulting diagram is readable.




Figure 10.1. Cohort-pathway editor. The design rail on the left holds Target Cohorts, Event Cohorts, and the Analysis settings (Collapse window, Minimum cell count, Maximum path length, Allow repeats). The canvas in the centre shows the data-source run table above the sunburst; the cohort coverage and legend live in the right rail.

10.3 Reading the result

The sunburst shows the index event at the centre and each subsequent step as a ring expanding outwards. The wedges in each ring are colour-coded by event cohort, drawn from a 20-colour palette — the legend table beside the chart maps each colour to the event cohort it represents, and the same colour appears at every step that event takes part in. The angular extent of every wedge is proportional to the number of persons whose pathway took that sequence.

Hovering a wedge shows the count at that step. Clicking a wedge opens the *path details* table beside the diagram, which lists every person-pathway containing that node, with columns for **Event**, **Remain**, **%**, **Diff**, and **%** — useful when you want to read off a specific path verbatim rather than guess from the geometry. A **Visualization** / **Tabular**

toggle on the canvas toolbar switches between the sunburst and a table-only view of the same pathways.

Generations are dispatched per source from the data-source run table at the top of the canvas: each row's primary action button cycles **Generate** → **Cancel** → **Rerun** as the run progresses. Beside it, a history icon (with a count badge once multiple runs exist) opens the **Previous runs** dialog scoped to that source, where you can reopen any historical diagram — the same dialog that ships with characterizations and incidence rates.

A diagram that looks visually busy but where the largest sequences only account for a few percent of the cohort suggests the event cohorts are too granular. Consolidate related medications into class-level event cohorts and re-run.

Incidence rates

The *incidence rate* module estimates how often new outcomes occur inside a target population per unit of person-time. For each (target cohort $\times$ outcome cohort $\times$ source) combination, ATLAS v3.0 returns persons-at-risk, time-at-risk, outcome cases, and the rate expressed per a chosen multiplier (commonly 1,000 person-years).

Note

The formal definitions of *incidence proportion* (per 1,000 persons) and *incidence rate* (per 1,000 person-years), and the "first event only, prior cases excluded" semantics ATLAS uses, are covered in *The Book of OHDSI*, chapter 11 [20].

11.1 Setting up an incidence rate analysis

Open Analysis and switch to the **Incidence Rates** tab of the Analysis Hub (Chapter 8), then click **New Incidence Rate**. Inputs:

Target cohorts. The populations whose outcome incidence you want to estimate.

Outcome cohorts. One per outcome you are studying.

Time at risk. Set start and end as anchor + offset pairs: each side picks an anchor date (**Cohort start** or **Cohort end**) and a number of days to offset. So an intent-to-treat 365-day window is *Start: Cohort start + 0 days, End: Cohort start + 365 days*; a per-protocol window pegs to **Cohort end**. Time-at-risk is automatically clipped at end of observation and at the first outcome.

Study window (optional). Restrict the analysis to a calendar date range with explicit **Start date** and **End date** fields. Use it for cohorts that should be analysed within a specific reporting period.

Stratify rules (optional). User-defined strata, each one a name plus a small criterion expression that mirrors the cohort-builder inclusion rules (concept set, attribute filters, temporal window). Add a rule per age band, per gender, per calendar year, or per custom partition; the result table breaks the rate down by every rule that evaluates to true for a person.




Figure 11.1. Incidence-rate editor. The design rail on the left holds Target cohorts, Outcome cohorts, Time at risk (Cohort start / Cohort end + offset), Study window, and Stratify rules. The canvas on the right shows the data-source run table above the result with the Target / Outcome / multiplier toolbar chips and the Treemap / Table toggle.

The time-at-risk window is the most consequential setting in this editor. Different choices (intent-to-treat vs per-protocol, etc.) can swing the rate by a factor of two or more on the same cohort. Report the TAR start and end with your rate and run a sensitivity check; the trade-offs are covered in *The Book of OHDSI* [20].

11.2 Reading the result

The result opens as a three-pane workbench: a collapsible design rail on the left (the same inputs you set when building the analysis), the centre canvas, and a right **Insights** rail.

The canvas stacks its elements top to bottom — the data-source run table, a toolbar, a one-line run-meta summary of the selected run, an **Outcome Overview** table, and finally the rates view itself. The **Outcome Overview** is a clickable Outcome / Persons / Cases table: selecting a row switches the rest of the canvas and the Insights rail to that outcome, which is how you move between outcomes when the analysis carries more than one. A **Treemap / Table** toggle on the toolbar switches the rates view between the treemap and the stratified table.

The toolbar also carries **Target**, **Outcome**, and **Multiplier** selectors and an **Export** menu (CSV export is implemented; SVG and PNG are listed as options).

The rates table columns are:

- **Stratum** — **Summary** for the headline row, then one row per stratify rule.
- **Persons** — persons at risk in that stratum.
- **Cases** — outcome occurrences inside the time-at-risk window.
- **TAR (years)** — total time-at-risk in person-years.
- **Rate / multiplier** — the incidence rate, expressed per the multiplier picked from the toolbar (the menu offers $\times 1$, $\times 10$, $\times 100$, $\times 1,000$, $\times 10,000$ and $\times 100,000$).

The table does not carry a confidence-interval column.

The treemap view shows nested rectangles whose area is proportional to the persons-at-risk in each stratum and whose fill follows the incidence-rate colour scale — the fastest way to spot strata where the rate is unusually high or low without scanning the table row by row.

11.2.1 The Insights rail

The right **Insights** rail summarises the selected outcome at a glance. It leads with a row of KPI cards — **Persons at risk**, **Outcome cases** (with the case proportion as a percentage), **Person-years**, and **Incidence rate** per the chosen multiplier — followed by a Cases-vs-Non-Cases donut and a stratification mini-table with **Stratum**, **Cases**, and **Rate** columns.

The canvas shows a data-source run table at the top with one row per connected source. Each row carries status, last-run timestamp, duration, and a single primary action button that cycles **Generate** → **Cancel** → **Rerun** as the run progresses, plus a history icon that opens the **Previous runs** dialog scoped to that source — the same dialog used by characterizations and pathways.

A high rate computed from a tiny denominator looks impressive in isolation. Always look at *persons at risk* alongside the rate; suppress or flag rows where it is below a sensible cell-count threshold.




Figure 11.2. The incidence-rate Insights rail — KPI cards for Persons at risk, Outcome cases with case proportion, Person-years, and Incidence rate; a proportion donut; and a Stratum / Cases / Rate mini-table.

Profiles

The **Profiles** area lets you look at one specific person's longitudinal record on one specific data source. It is useful for sanity-checking a cohort definition (does my cohort really catch the right people?) and for clinical review of individual cases when the data agreement permits.

12.1 Finding a person

1. Profiles is *not* in the top navigation. Reach a profile by following a **View profile** action from elsewhere in the application, or by opening a profile URL directly (Chapter 8 §8.3).
2. Pick a data source from the dropdown in the page actions, or pass it in the URL as `/profiles/{sourceKey}`.
3. Set the person ID in the inline-editable hero title at the top of the page: edit it in place and commit with Enter or by clicking away (blur), and use the reload icon beside it to refetch the record. The URL becomes `/profiles/{sourceKey}/{personId}` and the timeline loads.
4. Optionally append a cohort ID (`/profiles/{sourceKey}/{personId}/{cohortId}`) — the timeline will then highlight that cohort's membership window for the person.

12.2 The timeline view

The main view is a horizontal time axis spanning the person's observation period, with one swim-lane per OMOP domain. The twelve lanes ATLAS plots are: drug, condition, visit, procedure, observation, measurement, device, specimen, condition era, drug era, dose era, and death. The x-axis is labelled **Day (0 = cohort entry)**: every record is plotted at its day index, with day 0 — cohort entry — marked by a dashed *Cohort entry* markline. The axis is still a day index, not a calendar date.

Each lane is a sequence of marks at the day of each record in that domain; hovering a mark reveals the concept name and the day number. A row of filter chips above the chart lets you hide domains, and a brush on the timeline narrows the view to a day range; the active range appears as a closable chip ("Day *a* → Day *b*"). A separate events table below lists every record in detail, with **Concept Id**, **Concept Name**, **Domain**, **Start Day**, and **End Day** columns — day numbers rather than source codes or calendar dates — plus a **Search events** box to filter it.




Figure 12.1. Profiles landing page. Pick a source from the dropdown top-right, then type a person ID (or arrive via a deep-linked URL) to load the timeline.

For phenotype validation, set the cohort context before opening the profile (e.g. via the URL form `/profiles/{sourceKey}/{personId}/{cohortId}`). The cohort identifier appears as a closable chip in the page header so you know which cohort the records are being read against.

12.3 The rest of the profile page

12.3.1 The Demographics strip

A **Demographics** strip sits at the top of the profile, summarising the person in four fields: **Gender**, **Birth Year**, **Age-at-index**, and an **Events** count of the records plotted on the timeline.




Figure 12.2. The Demographics strip — Gender, Birth Year, Age-at-index, and Events count — across the top of the profile page.

12.3.2 The Observation periods band

Below the demographics is the **Observation periods** band: a thin, full-width strip with one rectangle per observation period, so a person with a gap in coverage shows two or more separated rectangles at a glance.

[screenshot]

The Observation periods band — a thin full-width strip with one rectangle per observation period.

Placeholder — to be replaced with screenshot of: The Observation periods band — a thin full-width strip with one rectangle per observation period.

12.3.3 The Highlights panel

The **Highlights** panel is a right-hand aside with two tabs, **Concepts** and **Sets**.

The **Concepts** tab lets you search the person's own concepts and assign each match one of six highlight colours; the matching marks recolour on the timeline so a chosen concept stands out against the rest. A **Clear all highlights** action resets every assignment.

The **Sets** tab is a disabled stub labelled "Concept-set highlighting coming soon" — concept-set-based highlighting is not yet available.

12.4 Validating a phenotype using Profiles

Profiles is the lightweight half of phenotype validation. The typical workflow is:

1. Generate the cohort against the source you want to validate on (Chapter 7).
2. Pick a small set of person IDs from the cohort — either by querying the cohort table directly, exporting the cohort result, or asking your administrator for a representative sample. Around a hundred persons is enough to find gross errors.
3. For each sampled person, open their profile with the cohort context set (`/profiles/{sourceKey}/{personId}/{cohortId}`). Read the records in the baseline window: do they look like patients with the condition? Read the records around the index date: is the entry event clinically plausible?
4. Tally what you see. Any obvious mismatches (a "type 2 diabetes" cohort that includes people with no glucose-lowering therapy and no diagnostic codes outside the index event) is a signal the rule is too loose; any common false negatives (clear diabetic patients excluded by an over-tight rule) is a signal it is too strict.
5. Adjust the cohort definition, save a new version (Chapter 14), regenerate, and resample.




Figure 12.3. The Highlights panel with its Concepts and Sets tabs; the Concepts tab shows a search box, colour swatches, and a Clear all highlights action.

Note

A sample of 100 charts is enough to find gross errors in a phenotype but not enough to estimate sensitivity, specificity, or positive predictive value with useful precision. For that, see the PheValuator section below.

12.5 Going further: PheValuator

ATLAS v3.0 itself does not formally evaluate a phenotype. For quantitative sensitivity, specificity, and PPV against a probabilistic gold standard, OHDSI ships *PheValuator* [18, 37, 38] as part of the HADES R suite [23]. The workflow is broadly: build the candidate phenotype and a tightly-specified "xSpec" cohort in ATLAS, export both as JSON, and run PheValuator in R against your CDM database — it fits a predictive model that yields per-source operating characteristics with confidence intervals.

The methodology, including how to construct the xSpec cohort, how to size the evaluation cohort, what target metrics make a phenotype publishable, and how PheValuator complements the CohortDiagnostics [28, 30] package, is covered in the package vignette and in *The Book of OHDSI*, chapter 16 "Clinical Validity" [22]. This manual documents what ATLAS provides; the evaluation step lives in R and outside the scope of the UI.

Feature analyses

A *feature analysis* is a reusable definition of a covariate or summary statistic. Characterizations (Chapter 9) pull their content from feature analyses, and once a feature is defined once it can be reused across studies — which is what makes characterizations comparable to each other.

13.1 Built-in versus custom features

ATLAS v3.0 surfaces the feature analyses defined in WebAPI — including the standard library that corresponds to the OHDSI `FeatureExtraction` [15] R package (part of the HADES suite [23]): demographic features, condition occurrence in fixed windows, drug exposure in fixed windows, visit counts, Charlson index components, and so on. Most users start by combining these in characterizations rather than authoring new ones.

13.2 The three feature-analysis types

To author a feature analysis, open `Analysis` and switch to the **Feature Analyses** tab of the Analysis Hub — the hub's default first tab (Chapter 8). **New Feature Analysis** opens the editor full-screen, where it asks for three choices side-by-side: **Analysis type** (how the feature is expressed), **Domain** (which CDM domain the rows come from), and **Stat type** (how the results are summarised).

The three analysis types are:

Preset (FeatureExtraction settings) a JSON blob of `FeatureExtraction CovariateSettings` — the most common path for importing a community-standard feature set without re-deriving it. The Design block on this type carries two helpers, **Load default covariate settings** and **Load default temporal covariate settings**, that pre-fill the JSON; useful when you want exactly what `createDefaultCovariateSettings(...)` would compute in R.

Criteria Set a small cohort-style criterion (concept sets, attribute filters, cardinality, temporal window). This is the type you reach for when the standard library does not cover a phenotype-specific feature you need to compute alongside the others. The visual criteria builder is not yet wired in v3.0; current authoring is through two raw JSON textareas (concept sets array and criteria group).

Custom SQL a raw SQL template parameterised against the results schema. Reserved for features that cannot be expressed as a criteria set.

The two statistical types are picked separately:

Prevalence the result is a binary covariate per person (present / absent), shown as a percentage with a count column on the characterization result page.

Distribution the result is a continuous covariate, summarised as mean, standard deviation, median, and quartiles.

Note

A typical baseline bundle combines the two stat types: demographics (age as distribution, gender as prevalence), condition / drug / procedure occurrences in the four standard time windows (prevalence), visit and observation counts (distribution), and Charlson comorbidity index components. This default is what most published OHDSI studies use unless the question specifically calls for something narrower.

The screenshot shows the 'Edit Feature Analysis' interface in ATLAS. At the top, there's a navigation bar with 'Data Sources', 'Concepts', 'Cohort Definitions', and 'Analysis'. The main form has a 'Name' field with 'Condition Group Era Start Long Term', a 'Description' field with 'One covariate per condition era rolled up to groups in the condition_era table starting in the long term window.', and two dropdown menus for 'Analysis type' (set to 'Preset') and 'Domain' (set to 'CONDITION'). Below these are two buttons: 'Load default covariate settings' and 'Load default temporal covariate settings'. At the bottom, there's a 'Design' section with a text area containing 'ConditionGroupEraStartLongTerm'.

Figure 13.1. Feature Analysis editor. **Analysis type** (Preset / Criteria Set / Custom SQL) and **Stat type** (Prevalence / Distribution) sit alongside **Domain**; the Design block carries the JSON or SQL body appropriate to the chosen type.

13.3 Reuse

Once a feature analysis is saved, every characterization that references it picks up the latest definition automatically. Unlike cohorts and concept sets, feature analyses do not yet carry their own version history or per-feature tagging in the UI — the data model has those fields but the editor does not expose them. Treat each save as the canonical state, and use characterization-level versioning (Chapter 14) to pin a feature analysis as it appeared on a given run.

PART V

Operations

Versions

Cohort definitions, concept sets, characterizations, pathways and incidence rates are all versioned. Every save creates a new version; older versions remain readable, and you can compare them, copy them, or restore them as the current working version. Feature analyses do not currently carry version history; their definitions are referenced directly by the characterizations that use them.

14.1 Where versions live

Each editable artefact carries a history icon (MDI-HISTORY) in its toolbar that opens a **Versions** dialog. The dialog lists versions in reverse chronological order; each row shows the version number, author, save timestamp, and a short comment if one was supplied, plus three actions: **Preview**, **Add comment** (or **Edit comment** when one is already set), and **Create a copy**. A small badge on the toolbar icon shows the current version count at a glance.




Figure 14.1. Versions dialog. Each historical version carries an author, timestamp, and optional comment, with **Preview**, **Add**, and **Create a copy** actions.

14.2 Previewing a version

Click a row in the Versions dialog to enter *preview mode*. The URL changes to e.g. `/cohortdefinition/{id}/version/{N}` and the editor renders the historical state. Preview mode is read-only — you cannot accidentally overwrite the past — and a banner across the top of the editor makes it obvious you are not on the current version.

To leave preview mode, click **Back to current version** on the banner, or navigate to `.../version/current`, which returns you to the live working copy.

14.3 Restoring a version

From a preview, hit **Save** on the editor toolbar; a confirmation dialog titled **Save as current version** appears and, on confirm, writes the previewed state back as a new version on top of the current one. The historical row is preserved; you have not deleted history, you have appended a new entry that re-asserts the older content.

Note

The save-as-current restore is currently wired for cohorts, concept sets, and characterizations. For pathway and incidence- rate previews the editor is read-only: you can open a historical version to inspect it, but you have to roll the change forward manually rather than restoring with one click.

Note

Restoring is non-destructive on purpose. If you decide the restore was a mistake, the version immediately before it is still in the list.

14.4 Pinning a published study

When a cohort, concept set, or characterization is referenced in a published paper or a regulatory submission, pin the exact version in the citation: cohort definition ID and version number, not just "the latest version of cohort 482". This is the cheapest way to guarantee a reader on a different ATLAS instance — or on the same one a year later — can recreate the exact analytical artefact. The version-preview URLs (`/cohortdefinition/{id}/version/{N}`) are stable; copy them straight into the paper's supplementary materials.

Note

When the cohort itself was reused from the OHDSI Phenotype Library, pin *two* version numbers: the local cohort version and the Phenotype Library release DOI it was imported from. Definitions accepted into a Library release are immutable for that release, but the same definition can be marked deprecated [D] or in error [E] in a later release; quoting the release DOI insulates your study from those later changes [26].

Administration

This chapter is for users with administrative permissions. Most researchers can skip it; come back when you need to grant a colleague write access, manage tags across the library, or check on long-running jobs.

15.1 The configuration panel

Administrative controls live in a side panel that opens from the **cog icon** in the top right of the navigation bar. The cog itself is gated: it is hidden entirely from users who hold none of the admin permissions, so a researcher account does not see it at all. Once open, the panel lists up to five tabs, each gated by a specific permission — a user with only one admin permission only sees one tab:

Cache (admin:cache). Patient and TrexSQL caches that per-source reports rely on. Trigger rebuilds, watch the progress, and review the timestamp of the last successful build. See "Caches and TrexSQL" below for what each cache does.

Data Sources (admin:source). The CDM databases WebAPI is connected to, listed with their source key, dialect, and connection status. This view is *read-only* — adding a source, changing a connection string, or rotating credentials happens against WebAPI's source tables directly (or via the Broadsea / atlasdb seed scripts).

Tags (admin:tags). Create, rename and merge the tags used across cohorts, concept sets and analyses.

Permissions (admin:security). The entry point for role and permission management. Selecting a role here opens the dedicated role page (covered below).

Jobs (job:*:get). The status view for long-running tasks — cohort generations, characterization runs, incidence-rate computations, pathway analyses, cache builds.

If a user with the cog visible somehow loses every admin permission, the panel renders a single info message — "You don't have access to any administrative settings." — in place of the tabs.

Note

The configuration panel is a side overlay rather than a separate page, so you can quickly check job status or cache state without leaving the page you were on.




Figure 15.1. Configuration panel on the **Cache** tab. Cache management lives at the top; the Patient Cache section below has a per-source row with the cache state (here *Not Built*) and a **Build Cache** action.

15.2 Caches and TrexSQL

Several ATLAS pages back queries that are expensive to recompute on every visit — vocabulary searches, source-level descriptive reports, cohort report rollups. WebAPI handles those through two layers of cache that the configuration panel’s **Cache management** tab exposes:

Patient cache. A per-source pre-aggregation that descriptive reports (Chapter 4, cohort **Reports** tab in Chapter 7) read from, and that the cohort builder uses to show *live patient counts* as a researcher edits a definition — saving a full **Generate** round-trip on every change. The cache sits in the source’s results schema. Initial build is a long-running job; once built, it is reused until the source data is refreshed.

TrexSQL cache. A query-level cache used by WebAPI’s SQL execution layer. Repeated vocabulary queries and concept-set lookups hit the cache rather than the database, so the second time a user opens the same concept set the response is essentially instant. The cache lives on disk on the WebAPI host (path set by the WebAPI environment variable `TREXSQL_CACHE_PATH`) and is shared across users.

Rebuild a cache after:

- the underlying CDM source has been refreshed (so the patient cache reflects the new data),
- the OHDSI vocabularies on the source have been updated (so the TrexSQL cache does not return stale concept information), or

- the cache was built against an older WebAPI version whose query shape has changed.

Cache rebuilds appear as jobs in the **Jobs** tab below; they can take from minutes to hours depending on the size of the source.

15.3 Roles and permissions

ATLAS v3.0 uses role-based access control. A *role* is a named bundle of *permissions*; a user is a member of zero or more roles. Open the configuration panel and pick **Permissions**, or navigate directly to `/config/roles` to see the role list.

Note

A deployment-wide runtime flag, `enablePermissionManagement` in `config-local.json` (Appendix A), can hide the permission-management UI for *everyone* regardless of who holds `admin:security`. When it is set to `false` the **Permissions** entry point is suppressed at the deployment level; the underlying RBAC model and permission strings described below are unchanged. If an administrator cannot reach the role pages even with the security permission, check this flag before adjusting roles.




Figure 15.2. Roles list inside the **Permissions** tab of the configuration panel. Click any role to open the role-detail page (**Users** / **Permissions** / **Utilities** tabs); **New Role** starts a fresh empty role.

15.3.1 Built-in roles

A typical deployment ships a few built-in roles that you should not delete: `public` (default for unauthenticated visitors where allowed), `user` (read access to most arte-

facts), and an admin-style role with full write access. Your organisation may have added project-specific roles such as a per-study analyst role.

15.3.2 Editing a role

Click a role to open its detail page at `/config/roles/{id}`. The page lets you edit the role's name and description inline at the top, then exposes three tabs in this order (**Users** is the default):

Users The users currently in the role. Add or remove members here. Changes take effect at the user's next request, usually immediately.

Permissions The flags that decide what the role's members can do — view cohorts, write cohorts, generate against specific sources, manage roles, and so on. Most permissions are per-source, which is how you give a contractor access to one database without exposing the rest.

Utilities Import, export and clone operations on the role itself — useful when you want to copy a vetted role bundle to a sibling deployment.

15.3.3 Creating, importing and deleting a role

The role list carries three companion dialogs:

New role opens a small dialog with a name and description; the new role starts empty and you flesh it out from the detail page above. Pick a clear, hyphenated name (e.g. `cardiology-analyst`). Role names are visible to all users, so avoid embarrassing or project-internal language.

Import reads a role JSON file produced by another ATLAS instance's **Utilities** tab. It is the right way to replicate a vetted role bundle from one deployment to another.

Delete removes the role after a confirmation dialog. Built-in roles are typically protected; only delete a role you own.

Removing a permission from a role takes effect immediately for every member of that role. If a long-running cohort generation is in flight and the role member loses generation permission mid-flight, the run will fail. Coordinate sensitive role changes with your users.

15.3.4 Permission reference

The permission strings ATLAS checks fall into four families. The table below lists the ones the frontend reads when deciding what to render.

Resource verbs `read:{resource}`, `write:{resource}`, `create:{resource}`, where *resource* is one of `cohort-definition`, `conceptset`, `cohort-characterization`, `feature-analysis`, `pathway`, or `incidence`.

Read gates the list and editor; write gates **Save** and **Delete**; create gates the **New ...** button on the list page.

Per-entity grants alongside the global verbs, WebAPI returns a per-entity record for each artefact the user has been granted access to individually: an `accessTypes` list (containing `READ` and/or `WRITE`) plus a separate `isOwner` boolean. Ownership implies both read and write; an explicit `WRITE` entry overrides a missing global `write`: permission. This is how you give a researcher full control over their own cohorts without granting blanket `write:cohort-definition`.

Admin `admin:cache` (Cache tab and patient-cache rebuilds), `admin:source` (Data Sources tab), `admin:tags` (Tags tab and tag-management actions on entities), `admin:security` (Permissions tab, role CRUD; also acts as a global read+write override on every artefact). Holding any of these makes the cog icon visible in the navbar. Source access stays gated by `admin:source` / `source:*:access`, not by `admin:security`.

Source access `source:*:access` gates which CDM sources a user can run cohort generations and reports against. Per-source variants (`source:<key>:access`) can grant access to a single source without the wildcard.

Jobs `job:*:get` gates the Jobs tab in the configuration panel.

Note

A common starter-role bundle for an analyst is `read:*`, `write:*` on the artefact resources plus `create:*`, `source:*:access`, and `job:*:get` — everything except `admin:*`. A read-only viewer role drops the `write:*` and `create:*` entries. Per-entity grants then layer on top to share specific cohorts with collaborators outside the analyst pool.

15.4 Jobs and long-running tasks

The **Jobs** section in the configuration panel is the status view for every long-running task in the system: cohort generations, characterization runs, incidence-rate computations, pathway analyses, and the per-source caches. Each row shows the job name, target source, status (**RUNNING**, **COMPLETED**, **FAILED**), submitted-by user, start and end times, and duration.

Refresh re-fetches the list from WebAPI — useful when you are watching a long generation finish.

Filtering by status lets you focus on running jobs or recently failed ones.

Clicking a row opens the job's log, where errors and SQL warnings are reported verbatim.

The Jobs section is the right place to diagnose "why is my generation taking so long" and "did my characterization actually finish". For administrators, scanning the failed-jobs

filter periodically catches WebAPI-side problems (connection timeouts, out-of-memory errors) before users notice them.

15.5 Audit information

ATLAS v3.0 stores the author and last modifier on every artefact and shows both in the list views. Together with the version history (Chapter 14) this is the basic audit trail. Deeper audit logs — of generation runs and login events — live in the WebAPI database and are typically inspected directly there rather than through the UI.

Plugins

ATLAS v3.0 supports plugins — independently developed front-end modules that appear inside the application as additional navigation entries and pages. From a user's perspective a plugin behaves like any other area of ATLAS v3.0; from an administrator's perspective it is a separately released JavaScript bundle that the host loads at runtime.

16.1 What plugins look like

When your deployment has plugins installed, they appear in the top navigation alongside the built-in entries. The content area below the nav is replaced with the plugin's own UI when you select it. Plugins share the active session, so you do not log in again, and they receive the currently selected data source automatically.



Figure 16.1. Top navigation on a deployment without plugins; loaded plugins would appear as additional entries to the right of **Profiles** (or replace built-in entries when configured to).

16.2 Common plugin categories

A given deployment may ship none, some, or all of these:

- **Custom analytics.** Statistical workflows specific to a research network, exposed as a plugin so they live alongside ATLAS v3.0's built-in analyses without forking the application.
- **Connectors.** Bridges to external systems — a redcap export, a SAS dataset writer, a registry submission tool.
- **Visualisations.** Dashboards or chart libraries that consume cohort or characterization output.

16.3 Trust and isolation

Each plugin runs in its own sandbox with a scoped message bus to the host. Plugins receive only the auth context and the API surface the host explicitly exposes; they cannot reach into stores they have not been granted. From a research integrity standpoint this matters, because a plugin that produced study results becomes part of the chain of custody — knowing exactly what version of which plugin produced a chart is something your audit process should care about.

Note

There is no in-application admin UI for installed plugins; the plugin manifest is a static `src/config/plugins.json` on the host, set at build or deploy time. If you are reviewing a publication that used ATLAS v3.0, asking the deployment operator for that file — and the build versions of the plugins it references — is a fair request.

16.4 The `plugins.json` manifest

The host loads its plugin manifest from `/config/plugins.json` at startup; the file is validated with a Zod schema (`PluginManifestSchema` in `src/models/PluginModels.ts`). A minimal manifest looks like this:

```
{
  "version": "1.0",
  "plugins": [
    {
      "id": "hello-world-plugin",
      "name": "Hello World",
      "version": "1.0.0",
      "entryPoint": "hello-world-plugin/index.js",
      "menuItems": [
        {
          "id": "main",
          "name": "Hello World",
          "route": "/plugins/hello-world-plugin/main",
          "icon": "mdi-hand-wave",
          "order": 0
        }
      ],
      "metadata": {
        "author": "Atlas Team",
        "description": "Demo plugin",
        "homepage": "https://github.com/ohdsi/atlas"
      }
    }
  ]
}
```

The relevant fields:

version manifest format version (currently "1.0").

plugins[] one entry per plugin. Each plugin specifies an `id` (unique within the manifest), a human-readable `name`, the plugin's `version`, and an `entryPoint`. A relative `entryPoint` is resolved against a hard-coded `<BASE_URL>/plugins/` —

the deployment's base path followed by `plugins/`. An `entryPoint` that is already absolute (starts with `/`) or a full http URL is used verbatim, which is how a plugin bundle hosted on a separate origin is referenced. The plugin's `menuItems[].route` **must** start with `/plugins/<id>/` — the host rejects the manifest otherwise.

fabMounts optional. Instead of — or in addition to — a nav entry, a plugin can mount as a *floating action button*. Each entry carries an `id`, a `label`, an `icon`, an optional `color`, and an optional `position` (`bottom-right`, `bottom-left`, `top-right` or `top-left`; the default is `bottom-right`). This lets a plugin surface a persistent action from anywhere in the app without taking a slot in the top navigation.

metadata optional descriptive fields; surfaced nowhere user-facing but useful for audit queries. `metadata.homepage` is validated as a URL and `metadata.icon` as a string; an invalid URL rejects the whole manifest.

The manifest also has an optional `settings` block that tunes the host without touching code:

enableHotReload reload plugins on bundle changes (development only; default `true`).

loadTimeout milliseconds to wait for a plugin bundle before failing the load (default 30000).

pluginsPath parsed and validated into the manifest, but **not currently honoured** by the plugin loader: relative `entryPoints` are always resolved against the hard-coded `<BASE_URL>/plugins/` described above, regardless of this setting. Treat it as reserved for now.

showLoadingIndicators whether the host shows a spinner while a plugin's bundle is fetching.

navigation.enabledCoreItems / **disabledCoreItems**: explicit lists used to hide built-in nav entries on a deployment that does not need them. `disabledCoreItems` takes precedence.

theme.primaryColor, **theme.logoUrl**, **theme.logoNavigateTo** colour and logo overrides. `theme.primaryColor` is regex-validated as a hex code (e.g. `#1f425a`); an unrecognised value rejects the manifest.

header.show... toggles for the navbar slots: `showNavBar`, `showFeedbackButton`, `showLanguageSelector`, `showConfigButton`, `showUserMenu`; plus `header.feedbackUrl` for the external feedback link target.

A common deployment hygiene check is "does the manifest only list plugins we expect to ship?" — since the manifest is checked into deployment config, a code review on `plugins.json` catches the same supply-chain risks as a code review on the application itself.

16.5 Errors and troubleshooting

If a plugin fails to load, the host shows a placeholder card in place of the plugin's UI with the failure reason. The most common causes are network errors fetching the bundle, version skew between plugin and host, and configuration errors in `plugins.json`. Reload the page after the operator has redeployed; if the problem persists, open the browser developer console and forward the single-spa lifecycle errors to your administrator.

Troubleshooting

The most common problems users hit on ATLAS v3.0, and how to recognise them. Items are grouped by symptom rather than by chapter; if you know what is wrong already, jump to the relevant reference chapter for the long form.

17.1 I cannot log in

Symptom: the credentials I usually use are rejected. Check the WebAPI clock. ATLAS uses JWT tokens signed by WebAPI; if the WebAPI host's clock has drifted more than a minute or two relative to the browser, every freshly issued token looks already-expired and the login appears to fail (Chapter 2).

Symptom: my OAuth/SAML/OIDC button is missing. The login providers are configured on the WebAPI side. Ask your administrator whether `SECURITY_AUTH_*` for that provider is set to `true`; on a fresh deployment only DB-auth is enabled by default (Appendix A).

Symptom: I expected anonymous access but get a login prompt. At startup ATLAS calls `user/me`. Anonymous browsing works only when the WebAPI is actually configured to allow it — then that call returns 200 with an anonymous permission set. If it returns 401 instead (anonymous access not enabled on WebAPI), and the frontend's `userAuthenticationEnabled` flag is on, the router opens the login modal. Ask your administrator to enable anonymous access on WebAPI, or expect to sign in.

Symptom: an OIDC/SAML redirect returns but I am still not logged in. The provider bounced you back to the callback route (`/openid/callback` or `/saml/callback`) but no bearer token arrived with it, so ATLAS reports *"Authentication failed - no token received"*. This is a server-side handshake problem, not a browser one: the WebAPI callback did not set the token ATLAS looks for. Check the provider and WebAPI `SECURITY_AUTH_*` configuration (redirect URIs, client secret, callback URLs) with your administrator.

Symptom: I am repeatedly logged out mid-task. The token lifetime is shorter than your session. ATLAS shows an extend-session dialog before the token expires; if the page is not active the token quietly expires. Auto-saved cohort drafts are restored from session storage on next login, but unsaved analysis runs are lost.

17.2 I cannot see any data

Symptom: the Data Sources page lists my source but every report tab is empty. The OHDSI Achilles routine has not been run against that source's results schema. Atlas reads from the Achilles result tables; if they are missing, every tab renders blank even when the source itself is connected. This is the single most common "set up but nothing shows" cause. See Appendix A for how to run Achilles.

Symptom: the Data Sources page is empty. WebAPI has no sources registered. Check `<host>/WebAPI/source/sources` — if it returns an empty JSON array, no source rows exist. Add them via the WebAPI database tables or the atlasdb seed scripts (Appendix A).

Symptom: per-source record-count columns or live patient counts stay blank. Those numbers come from the per-source patient cache. If the cache has not been built (or is stale after a CDM refresh), they remain dashes. An administrator can rebuild the cache from the configuration panel's **Cache** tab; cache builds run as long-running jobs visible in the **Jobs** tab (Chapter 15).

17.3 Cohort generation does not produce what I expect

Symptom: the cohort generates zero rows. Click **Inclusion report** on the source's row in the **Generation** section. The first step of the attrition funnel is the entry-event count; if it is already zero, your concept set has no records on this source. Either the concept set resolves to nothing (open the set's **Included Concepts** tab and check the per-row record-count columns), or the entry-event window or attribute filters exclude every candidate. Run the cohort against a source you know has the relevant data first.

Symptom: the cohort produces vastly more rows than expected. Almost always one of: (i) **Cohort entry on** is set to **all** when you wanted **earliest event**, multiplying the row count; (ii) a concept set has **Descendants** ticked on a parent that pulls in unintended children (Chapter 6).

Symptom: an inclusion rule drops nobody. The rule is either trivially satisfied (cardinality too lax, temporal window too wide) or redundant with an earlier rule. Tighten or remove. The inclusion-rule debugger guidance in Chapter 7 walks through this.

Symptom: the same JSON produces wildly different counts on two sources. Almost always an ETL or vocabulary-mapping difference, not a real clinical signal. Run the same cohort on a third source (Chapter 7); if its attrition pattern matches one of the two but not the other, the outlier is the ETL-suspect source.

Symptom: generation takes much longer than expected. Open the configuration panel's **Jobs** tab (cog icon → **Jobs**) and inspect the running job's log. WebAPI prints SQL warnings and database-side errors verbatim. Common causes: a concept set with

thousands of resolved members hammering `CONDITION_OCCURRENCE`, or a missing index on the results schema (an administrator concern).

17.4 Concept set issues

Symptom: the search returns no results for a term I know is in the data. Check the Vocabulary filter — if it excludes the vocabulary your data uses you will see nothing. ICD codes live in the ICD vocabulary; SNOMED standards live in SNOMED; non-US drugs live in RxNorm Extension rather than RxNorm. Clear the filter and re-search if uncertain.

Symptom: a concept set with Descendants ticked covers thousands of unrelated terms. Either the parent concept is too high in the hierarchy — **Descendants** pulls down everything — or the parent concept is itself non-specific. Move down the hierarchy or narrow the parent. The descendant record count (**DRC**) on the row in the **Included Concepts** tab is your in-place sanity check.

Symptom: my concept set looks right but the cohort matches nobody. The source's `*_concept_id` columns may be populated from a different vocabulary than the one you selected (e.g. the ETL maps to SNOMED while you searched on ICD). Toggle **Mapped** on the rows in the **Included Concepts** tab so source codes that map up to your standard concept also count, and confirm the record-count columns on the row become non-zero before re-generating.

17.5 Permissions: I cannot see or do something

Symptom: the cog icon for the configuration panel is not visible. The cog is hidden from users who hold none of the admin permissions (`admin:cache`, `admin:source`, `admin:tags`, `admin:security`, `job:*:get`). Ask your administrator to add you to a role that carries the relevant permission (Chapter 15).

Symptom: the configuration panel opens but only some tabs are present. Each admin tab is gated by its own permission, so a role with only `admin:cache` sees only the **Cache** tab. The other tabs are not disabled; they are not rendered at all. This is expected.

Symptom: Save or Delete is disabled on a cohort or concept set I am editing. You do not hold write permission on this entity. Either you are not the author and your role lacks the relevant `write:cohort-definition` / `write:conceptset` permission, or the entity is locked by another concurrent edit. A per-entity `WRITE` grant from your administrator overrides the missing global permission for a specific artefact.

17.6 Errors that show up in the browser console

CORS errors. The frontend is on a different origin from WebAPI, and WebAPI's `SECURITY_CORS_ALLOWED_ORIGINS` does not list the ATLAS origin. Add the origin or put both behind the same reverse proxy (Appendix A).

401 Unauthorized. The bearer token has expired or was rejected. Reload the page; ATLAS will prompt for a fresh login.

404 Not Found on a /WebAPI route. The WebAPI base the frontend calls points at the wrong place. ATLAS reads it at runtime from `api.url` in `config-local.json` (default `/WebAPI`, i.e. same-origin); the compiled-in `VITE_WEBAPI_URL` is only a fall-back. On the compose stack the same-origin `/WebAPI` path is served by Caddy's `/WebAPI reverse_proxy` handler, so a 404 there means either `api.url` is wrong or Caddy is not forwarding to the WebAPI service. The appendix's verify-step is to call `<host>/WebAPI/info` and check the JSON response (Appendix A).

500 Internal Server Error on a /WebAPI route. A WebAPI-side problem. The response body usually carries a human-readable error; if not, check the WebAPI log.

17.7 Plugin problems

Symptom: a plugin entry in the navigation shows an error card instead of its UI. The plugin bundle failed to load. Common causes: a network error fetching it, version skew between plugin and host, or a configuration error in `plugins.json` (Chapter 16). The browser developer console usually carries a single-spa lifecycle error that points at the issue.

17.8 When the answer is not in this manual

- Live WebAPI reference: `<host>/WebAPI/swagger-ui/` (covered in Appendix A).
- Atlas3 source / issue tracker: github.com/OHDSI/Atlas3.
- OHDSI community forum: forums.ohdsi.org.
- The Book of OHDSI [20] for the methodology behind any feature.

Setup guide: WebAPI and ATLAS v3.0

This appendix walks through a working installation of WebAPI and the ATLAS v3.0 frontend. The goal is to get a developer or evaluator from zero to a running web application that talks to a real PostgreSQL backing store. For production deployments, follow your organisation's infrastructure standards rather than the recipes here — this chapter is intentionally minimal.

ATLAS v3.0 requires **WebAPI v3.0**. The 2.x line of WebAPI that the production ATLAS uses — and that ships with the OHDSI Broadsea distribution — is *not* compatible. The v3.0 backend introduces schema and endpoint changes the frontend relies on, and the frontend assumes the v3.0 contract from its first request. The development docker compose in the Atlas3 repository builds its WebAPI image *on top of* the matching base image `ghcr.io/ohdsi/webapi:3.0-dev` for exactly this reason; if you build WebAPI from source instead, check out the `webapi-3.0` branch explicitly.

A.1 The pieces

A working ATLAS v3.0 deployment has three independent components:

A relational database holding the OHDSI WebAPI schema, the OMOP Common Data Model schemas you want to analyse, and any result schemas. PostgreSQL is the simplest choice and the one the development `docker-compose.yml` ships with; Microsoft SQL Server, Oracle, Amazon Redshift, and Snowflake are all supported by WebAPI.

The WebAPI service a Java/Spring application that exposes the REST endpoints ATLAS talks to, executes cohort SQL against the configured CDM sources, and runs Flyway migrations against its own schema on first boot.

The ATLAS v3.0 frontend a Vue 3 single-page application, served either by the Vite development server or as a static bundle behind a reverse proxy.

In a Docker Compose setup these are three services. In a production deployment they may be one database cluster, one Java application server, and a static-asset host.

A.2 Prerequisites

- Docker and Docker Compose v2 (the fastest path).
- *or* for native installs: PostgreSQL 12+, JDK 17, Maven 3.9+, Node.js 18+ and npm.
- Git.
- Roughly 4 GB of free disk for container images and the WebAPI Flyway baseline.
- A WebAPI build from the v3.0 line — container image `ghcr.io/ohdsi/webapi:3.0-dev` or a local build of the matching branch from the [WebAPI repository](#).

A.3 The Docker Compose path (recommended for evaluation)

The ATLAS v3.0 repository ships a `docker-compose.yml` that brings up Postgres, WebAPI, a database initialiser, and the frontend behind Caddy with a self-signed certificate. From a fresh checkout:

```
git clone https://github.com/OHDSI/Atlas3.git
cd Atlas3
docker compose up
```

The stack does not pull a prebuilt WebAPI image verbatim. The `atlas3-webapi` service *builds* a local image (`docker/webapi/Dockerfile`) that starts FROM `ghcr.io/ohdsi/webapi:3.0-dev` and overrides the bundled TrexSQL plugin jar with a pinned release (`TREXSQL_VERSION=0.2.3`, downloaded from [OHDSI/trex](#)). The database is the `ohdsi/broadsea-atlasdb:2.3.0` image (service `atlas3-postgres`, network hostname `broadsea-atlasdb`), and a separate `atlas3-db-init` step runs the `atlasdb/` seed scripts once WebAPI reports healthy. ATLAS v3.0 will not work against the WebAPI 2.x line, which is why the image is pinned to the v3.0 base. The first boot takes a few minutes while WebAPI runs its Flyway migrations into the `webapi` schema (baseline 3.0.0) and the init step seeds a default user.

When the stack is healthy:

- ATLAS v3.0 is at <https://localhost/atlas/> — Caddy redirects the bare `/` to `/atlas/` and mints a self-signed certificate, which you must accept. Routing is *hash-based*, so deep links and OAuth/OIDC/SAML callbacks carry the in-app path after a `#`, e.g. <https://localhost/atlas/#/cohortdefinitions>.
- WebAPI is at <http://localhost:8080/WebAPI/info>; that endpoint returns a JSON blob with the version once the service is up. Caddy also reverse-proxies it at <https://localhost/WebAPI/> on the same origin as ATLAS.
- PostgreSQL is *not* published to the host. It is reachable only on the compose network, as `broadsea-atlasdb:5432` (equivalently `atlas3-postgres:5432`). To reach it from the host, exec `psql` inside the container or add an explicit ports mapping to the service yourself.

Note

The Caddy front end reads three environment variables (defaults in parentheses): `CADDY_HOSTNAME` (`localhost`), the site address it answers on; `CADDY_TLS_DIRECTIVE` (`tls internal`), the TLS directive, whose default auto-mints a local certificate; and `WEBAPI_URL` (`https://localhost/WebAPI`), the WebAPI base advertised to the frontend. Set them in a `.env` file to serve on a real hostname with a real certificate.

The default credentials match the values seeded by the `atlasdb` scripts; check that folder for the actual username and password and *change them before any non-local use*.

The default `POSTGRES_PASSWORD` in the compose file is `mypass`. It is a placeholder for local development. Set `POSTGRES_PASSWORD` in a `.env` file (or your secrets manager) before running anywhere a hostile network can reach.

Note

The compose file already enables the TrexSQL query cache (`TREXSQL_ENABLED=true`, `TREXSQL_CACHE_PATH=/data/cache`, `HOME=/data`) and persists it on the `atlas3-webapi-data` named volume. Live patient counts in the cohort builder and the inclusion-rule rail only appear once TrexSQL is enabled *and* a per-source patient cache has been built — trigger the first build from Configuration → Cache management after login. See Chapter 15, "Caches and TrexSQL".

A.4 Native install

If you want each component on the host without Docker, the steps below mirror the compose recipe.

A.4.1 1. Database

1. Install PostgreSQL 12 or newer.
2. Create a database (`ohdsi` is a common choice) and a `webapi` schema.
3. Create a role with `CREATE` privileges on that schema — WebAPI needs to run Flyway migrations against it on startup.
4. Optionally, restore one or more OMOP CDM databases as additional schemas. These are the data sources the application will offer in **Data Sources**.

A.4.2 2. WebAPI v3.0

1. `git clone --branch webapi-3.0 https://github.com/OHDSI/WebAPI.git`. The default branch of the WebAPI repository tracks the production 2.x line, which ATLAS v3.0 cannot talk to; explicitly check out the `webapi-3.0` branch.

2. Create a Maven settings profile or `application.properties` override that maps the database connection details to the property keys WebAPI expects. The compose file lists the relevant variables verbatim — see `atlas3-webapi.environment` in `docker-compose.yml` for the canonical names (`DATASOURCE_URL`, `DATASOURCE_USERNAME`, `DATASOURCE_PASSWORD`, `DATASOURCE_OHDSI_SCHEMA`, the `SPRING_FLYWAY_*` keys, and the `SECURITY_AUTH_DB_DATASOURCE_*` keys for database-backed authentication). PostgreSQL is the default datasource dialect; no extra Maven profile is required for it. For Oracle, MSSQL, Redshift, Snowflake, BigQuery, etc. activate the matching `webapi-*` profile from `pom.xml` to pull in the JDBC driver.
3. Build the Spring Boot fat-jar:


```
mvn package -Dpackaging.type=jar
```

 The result is `target/WebAPI.jar`. (Without the `-Dpackaging.type=jar` flag, Maven produces a WAR; WebAPI v3.0 is designed to run as the fat-jar, not in an external servlet container.)
4. Run it with `java -jar target/WebAPI.jar`, passing the `DATASOURCE_*` and `SPRING_FLYWAY_*` settings as environment variables or via `--spring.config.location`. WebAPI will run Flyway migrations against the configured schema on first boot.
5. Confirm <http://localhost:8080/WebAPI/info> returns the version JSON.

A.4.3 2a. Enable the TrexSQL query cache

Live patient counts in the cohort builder, the inclusion-rule rail, and several descriptive views all read from *TrexSQL*, a query-result cache that **is not part of the WebAPI fat-jar** and must be added separately. The WebAPI image used by the Docker Compose stack already bakes it in — in fact the stack rebuilds that image to pin a specific TrexSQL release (see *The Docker Compose path* above) — so this step matters only for native installs that built `WebAPI.jar` themselves.

Only **one artefact** is required: the DuckDB-backed `trexsql.jar`, a self-contained uberjar published on the [OHDSI/trex](#) GitHub releases page. It embeds its native DuckDB library inside the jar, so the separate per-platform shared library (`libtrexsql.so` / `.dylib` / `.dll`) and the JNA platform-directory layout that older instructions described are **no longer needed**. Match the jar version to the WebAPI build you assembled — the `ARG TREXSQL_VERSION` line in the `docker/webapi/Dockerfile` of the Atlas3 repo pins the combination used for the official image and is the safest reference. Drop the jar into a `plugins` directory the launcher will scan:

```
/opt/webapi/plugins/trexsql.jar
```

Then set three environment variables and one JVM flag on the WebAPI process:

`TREXSQL_ENABLED=true` turns the cache on. With this unset or `false` the `/WebAPI/trexsql/*` endpoints respond 404, the cohort builder hides the patient-

count bar entirely, and the Configuration → Cache management page omits the TrexSQL section.

TREXSQL_CACHE_PATH absolute filesystem path where the cache files live, e.g. `/var/lib/webapi/cache`. WebAPI creates it on first boot and writes per-source subdirectories underneath. Place it on a volume that survives restarts — losing the directory forces a full rebuild for every source.

HOME writable directory the cache uses for per-user working files. The compose stack sets `HOME=/data` alongside `TREXSQL_CACHE_PATH=/data/cache`; mirror that pattern by pointing `HOME` at a writable location owned by the WebAPI process user. Without a writable `HOME`, cache initialisation fails on startup.

`-Dloader.path=/opt/webapi/plugins` JVM flag that tells Spring Boot's properties launcher to scan the plugins directory for additional jars. Without it, the `trexsql.jar` on disk is never loaded and the `TREXSQL_ENABLED` flag is a no-op.

Putting it together, the launch command becomes:

```
java
  -Dloader.path=/opt/webapi/plugins
  -jar target/WebAPI.jar
```

with `TREXSQL_ENABLED`, `TREXSQL_CACHE_PATH`, `HOME` and the usual `DATASOURCE_*` variables exported in the environment. After restart, confirm the cache endpoints are mounted by hitting one of them for a known source key (substitute `<sourceKey>` for a key returned by `/WebAPI/source/sources`):

```
curl -s \
  http://localhost:8080/WebAPI/trexsql/<sourceKey>/cache/status
```

should return a JSON status object (`not_built`, `building`, `ready`, `stale`, or `error`) rather than a 404. Then sign in to ATLAS v3.0 and open Configuration → Cache management; each registered source appears with a **Build** button. The first build runs as a long job (Chapter 15); only after it finishes does the cohort builder start showing live patient counts for that source.

Note

Per-user access to TrexSQL is gated by the `trexsql:*` permission group (Chapter 15). The built-in admin role has the full set; for non-admin users, add the relevant flags to their role before they can use live patient counts or trigger a rebuild.

A.4.4 3. Add data sources to WebAPI

WebAPI reads its source list from a database table populated by the `atlasdb` initialisation scripts (or by the [Broadsea](#) stack). Each row points WebAPI at one CDM database with a JDBC URL, schema names for CDM and results, and a friendly source key. Until at least one source exists, the ATLAS v3.0 **Data Sources** page is empty.

A.4.5 4. ATLAS v3.0 frontend

1. `git clone https://github.com/OHDSI/Atlas3.git && cd Atlas3.`
2. `npm install.`
3. Point the frontend at WebAPI and configure authentication. In ATLAS v3.0 this is done at *runtime*, not at build time: on startup the app fetches `config-local.json` from the same directory as `index.html`, Zod-validates it, and merges it over the built-in defaults (see *Externalized runtime configuration* below). The old build-time `VITE_AUTH_ENABLED`, `VITE_AUTH_SKIP_LOGIN` and `VITE_AUTH_PROVIDERS` variables are no longer read. For local development, drop a `config-local.json` into the served root (e.g. `public/`) with at least `api.url`, `userAuthenticationEnabled`, and your `authProviders`.
4. For development, run `npm run dev`; the dev server listens on `http://localhost:5173` and proxies `/WebAPI` requests to `http://localhost:8080` — a value hard-coded in `vite.config.ts`. To point the dev server at a WebAPI on a different host or port (a remote shared instance, a native build on 8081, etc.), edit the `proxy.target` in `vite.config.ts`. `VITE_WEBAPI_URL` still supplies the bundle's compiled-in default WebAPI base, but it is superseded at runtime by `api.url` from `config-local.json`; either way it does not affect the dev proxy target.
5. For production, run `npm run build` and serve the contents of `dist/` from a static host. A reverse proxy must forward `/WebAPI` to the Java service so the same origin serves both. The repository's Caddyfile shows a working configuration.

Note

`npm run dev` is the right entry point for trying things out or for plugin development; it runs the Vite dev server with hot module reload. The compose stack instead builds the static bundle into the Caddy image, which is what a production deployment would do.

A.4.6 Externalized runtime configuration

Auth and feature configuration is loaded at *runtime*, so the same static bundle can be re-pointed without a rebuild. On startup the app fetches `config-local.json` from the directory that serves `index.html`, validates it against a Zod schema, and merges the recognised keys over the built-in defaults (a shallow merge; arrays replace entirely). A missing or malformed file falls back to defaults with a console warning. In the Docker image the file is baked in at `/srv/atlas/config-local.json` (copied from `docker/config-local.json`).

The keys the loader recognises:

`api.url` WebAPI base URL (default `/WebAPI`, i.e. same-origin behind the reverse proxy).

Overrides the compiled-in `VITE_WEBAPI_URL` default.

userAuthenticationEnabled whether a session is required. The shipped image sets it true; left false (the default) the app runs against WebAPI's anonymous access (Chapter 2).

enableSkipLogin offers a "continue without signing in" path where the deployment allows it.

enablePermissionManagement hides the permission-management UI deployment-wide when false (Chapter 15).

authProviders the list of login providers (name, login URL, icon, whether it uses the credentials form) rendered in the login modal.

Feature flags **enablePersonCount**, **enableTaggingSection**, **defaultLocale** and **pollInterval** tune optional UI and polling behaviour.

A.5 First-time login and creating a user

The default DB-auth path expects a row in the `webapi.auth_user` table whose password column is a bcrypt hash of the chosen plaintext. The compose seed scripts in the `atlasdb/` folder create an initial administrator. To add more users, either insert rows into `webapi.auth_user` directly or, if your deployment exposes LDAP/OAuth/SAML/OIDC, configure the corresponding `SECURITY_AUTH_*` group of properties on the WebAPI side and let users log in via that flow instead.

After signing in, an administrator should:

- visit Configuration → Permissions and set up the role bundles their organisation needs (Chapter 15);
- confirm the **Data Sources** page lists each connected CDM and that its dashboard renders without errors (Chapter 4);
- run a small smoke-test cohort against each source to validate end-to-end execution (Chapter 7).

A.6 Common installation issues

The ATLAS page loads but every list is empty. WebAPI is reachable but no CDM sources are configured. Add rows to the WebAPI source tables (or run the `atlasdb` init scripts).

CORS errors in the browser console. The frontend and WebAPI are on different origins and WebAPI's `SECURITY_CORS_ALLOWED_ORIGINS` does not include the ATLAS v3.0 origin. Either add the origin to the allow-list or put both behind the same reverse proxy.

Flyway error on WebAPI startup. The configured database user lacks privileges on the `webapi` schema, or a previous WebAPI version left a migration in an inconsistent state. Inspect the WebAPI log; Flyway prints the failing migration explicitly.

Login fails with valid credentials. Check the WebAPI clock against the host running the browser. A skew larger than the JWT clock-tolerance window makes every freshly issued token look already-expired.

Data Sources page is empty for every tab. The most common case: the source is configured but *Achilles* [9] has not been run against it. ATLAS v3.0's descriptive reports read from the Achilles results tables in the source's results schema; if those tables are missing, every tab renders empty even though the source itself is reachable. See the next section.

A.7 Verifying a clean setup

After the stack is up, run these three checks before declaring victory:

1. `curl -sk https://localhost/WebAPI/info` should return JSON containing the WebAPI version and build date.
2. `curl -sk https://localhost/WebAPI/source/sources` should return a non-empty JSON array with one entry per configured CDM. An empty array means no sources are registered, and every ATLAS v3.0 list page that depends on a source will appear empty.
3. In the browser, open the developer console (Ctrl+Shift+I in Chromium) and watch the Network tab while loading **Data Sources**. You should see 200 responses for the `/source/sources` call and for each per-source report. Repeated 401 responses point at an authentication problem, 404 at routing, 500 at WebAPI itself — the response body usually has a clear error.

A.8 Achilles: the missing prerequisite for descriptive reports

OHDSI Achilles [9] is a separate R package that analyses a CDM database and writes summary tables into the same source's results schema. ATLAS v3.0's **Data Sources** reports *read from those Achilles tables*; if Achilles has never run, the source is connected but the dashboards are blank. Most evaluators hit this once and remember it forever.

For each source you want to use:

1. Install Achilles in an R environment that can reach the CDM:
`install.packages("Achilles")` (or via `remotes::install_github("OHDSI/Achilles")`).
2. Run `Achilles::achilles(connectionDetails, ...)` with the CDM and results schema names. On a multi-million person database this can take from minutes to hours.
3. Re-run Achilles after every CDM refresh; the results schema is rebuilt from scratch each time.

Two adjacent OHDSI tools you may want to install at the same time: the Data Quality Dashboard [14] for systematic data-quality checks against the CDM (Chapter 4), and PheValuator [18, 37] for formal phenotype evaluation (Chapter 12).

A.9 Related: Broadsea

If the recipe in this appendix feels too low-level, the OHDSI [Broadsea](#) [13] project bundles WebAPI, ATLAS, Achilles, the Data Quality Dashboard, and a sample CDM into a single Docker stack with a guided setup. It is the fastest way to demo the whole OHDSI toolchain on a laptop.

The current Broadsea release ships ATLAS 2.x and the matching WebAPI 2.x backend, neither of which works with ATLAS v3.0. Broadsea support for ATLAS v3.0 / WebAPI v3.0 is on the OHDSI roadmap and expected in an upcoming release; until that lands, swapping in the v3.0 components manually is required and non-trivial (the database seed scripts and the WebAPI image both need to change in lockstep). For a v3.0 stack today, use the `docker-compose.yml` that ships in the Atlas3 repository — it is the supported path and what the rest of this appendix documents.

A.10 Further reading and API reference

This manual documents what users do in the ATLAS v3.0 interface; it is not a reference for the WebAPI backend that the UI talks to. For power users who want to call WebAPI directly — to script cohort generation, batch concept-set imports, or build a plugin against the same endpoints the UI uses — the canonical references live upstream:

ATLAS v3.0 frontend repository github.com/OHDSI/Atlas3 — source, issue tracker, build instructions.

WebAPI repository github.com/OHDSI/WebAPI — the v3.0 backend; check out the `webapi-3.0` branch for the version this manual documents.

Live WebAPI reference `<host>/WebAPI/swagger-ui/` — the deployed instance's own swagger UI lists every endpoint with request and response shapes. On a local Docker Compose stack the URL is <http://localhost:8080/WebAPI/swagger-ui/>.

Book of OHDSI ohdsi.github.io/TheBookOfOhdsi — the conceptual reference for OMOP / ATLAS / the OHDSI toolchain. Cited throughout this manual.

The bibliography section at the back of the manual lists the papers and packages cited from individual chapters, including HADES [23], FeatureExtraction [15], Achilles [9], and PheValuator [37].

Glossary

CDM (Common Data Model) The OMOP Common Data Model is a person-centric relational schema that standardises observational health data so the same analytic code runs against any compliant dataset.

Cohort A set of persons who satisfy one or more inclusion criteria for a duration of time. In ATLAS v3.0 a cohort is defined by an entry event, optional inclusion rules, and an exit rule.

Concept A single entry in the OMOP standardised vocabularies, such as a SNOMED condition, an RxNorm ingredient, or a LOINC measurement.

Concept set A reusable, named collection of concepts, with options to include descendants, include mapped source codes, and exclude specific concepts.

Standard concept A concept the OHDSI vocabulary marks as canonical for its domain. Source codes (ICD-10, NDC, etc.) usually map up to a standard concept. Cohort logic should reference standard concepts so it is portable across data sources.

Era A continuous period of clinical state derived from underlying events — for example, a drug era is built from contiguous drug exposures with a configurable persistence gap.

Index date The cohort entry date for a person; most criteria are expressed as windows relative to the index date.

WebAPI The REST backend that ATLAS v3.0 talks to. WebAPI executes cohort generation, vocabulary search, and reporting against the underlying CDM databases.

Source / data source A configured CDM database. ATLAS v3.0 lists every source that WebAPI is connected to; reports and cohort generation always run against a chosen source.

Characterization An analysis that summarises clinical features of a target cohort, optionally compared to outcomes, using a set of feature analyses.

Feature analysis A reusable definition of a covariate or summary statistic (e.g. "drug exposure in the 365 days before index"). Used by characterizations.

Cohort pathway A visualisation of the order in which event cohorts occur for the persons in a target cohort.

Incidence rate The number of new outcome events per person-time-at-risk inside a target cohort.

Inclusion rule A criterion applied after the entry event that a person must satisfy to remain in the cohort.

Censoring event A clinical event that ends a person's cohort membership before the configured exit window would otherwise close it.

PHOEBE A vocabulary-recommendation engine packaged with WebAPI that suggests additional concepts likely to belong in a concept set, drawn from co-occurrence patterns across an OMOP network. ATLAS v3.0 surfaces it as the **Recommend** tab inside the concept-set editor.

TrexSQL A query-result cache shared across users on a WebAPI host, used to short-circuit repeated vocabulary and concept-set lookups. Distinct from the patient cache below.

Patient cache The per-source, pre-aggregated dataset that descriptive reports and the live patient-count bar in the cohort builder read from. Built on demand from the configuration panel's **Cache** tab; rebuilt when the underlying CDM or vocabulary refreshes.

Previous runs The historical generations of an analysis, listed in the **Previous runs** dialog. The dialog is scoped to a single data source and opened from the history icon (MDI-HISTORY, with a count badge once multiple runs exist) on each row of the data-source run table at the top of the Characterization, Incidence Rate, and Pathway editors.

Subgroup analysis An overlapping selection of the population inside a characterization, defined like a small inclusion rule. Each subgroup adds an extra column to the result. (Internally referred to as *strata*; the UI label is "Subgroup analyses".)

Sunburst The radial diagram used to visualise cohort pathways: the index event sits at the centre, each subsequent step is a ring expanding outwards, wedge area is proportional to person count.

Achilles The OHDSI R routine that computes the descriptive statistics ATLAS reads from for the **Data Sources** reports. A source is "Achilles-processed" when those routines have been run against its results schema.

Athena OHDSI's vocabulary repository (<https://athena.ohdsi.org>). The vocabulary tables Athena publishes are loaded into WebAPI; Atlas's vocabulary search runs against that data, so two ATLAS instances on the same release but with different Athena snapshots can return slightly different concept IDs.

HADES The OHDSI R-package suite (*Health Analytics Data to Evidence Suite*). HADES packages such as `FeatureExtraction`, `CohortDiagnostics` and `PheValuator` consume the JSON ATLAS exports for analyses ATLAS itself does not run.

Bibliography

- [1] Peter C. Austin. Balance diagnostics for comparing the distribution of baseline co-variates between treatment groups in propensity-score matched samples. *Statistics in Medicine*, 28(25):3083–3107, 2009.
- [2] EHDEN. Phenotype definition, characterisation and evaluation. EHDEN Academy course, https://www.ehden.eu/course_phenotypedefinitioncharacterisationevaluation/, 2024.
- [3] Health Data Research UK. HDR UK phenotype library. <https://phenotypes.healthdatagateway.org/>, 2025.
- [4] George Hripcsak, Jon D. Duke, Nigam H. Shah, Christian G. Reich, Vojtech Huser, Martijn J. Schuemie, Marc A. Suchard, Rae Woong Park, Ian Chi Kei Wong, Peter R. Rijnbeek, Johan van der Lei, Nicole Pratt, G. Niklas Norén, Yu-Chuan Li, Paul E. Stang, David Madigan, and Patrick B. Ryan. Observational health data sciences and informatics (OHDSI): Opportunities for observational researchers. *Studies in Health Technology and Informatics*, 216:574–578, 2015.
- [5] George Hripcsak, Patrick B. Ryan, Jon D. Duke, Nigam H. Shah, Rae Woong Park, Vojtech Huser, Marc A. Suchard, Martijn J. Schuemie, Frank J. DeFalco, Adler Perotte, Juan M. Banda, Christian G. Reich, Lisa M. Schilling, Michael E. Matheny, Daniella Meeker, Nicole Pratt, and David Madigan. Characterizing treatment pathways at scale using the OHDSI network. *Proceedings of the National Academy of Sciences*, 113(27):7329–7336, 2016.
- [6] George Hripcsak, Matthew E. Levine, Ning Shang, and Patrick B. Ryan. Effect of vocabulary mapping for conditions on phenotype cohorts. *Journal of the American Medical Informatics Association*, 25(12):1618–1625, 2018.
- [7] Michael G. Kahn, Tiffany J. Callahan, Juliana Barnard, Alan E. Bauck, Jeff Brown, Bruce N. Davidson, Hossein Estiri, Carsten Goerg, Erin Holve, Steven G. Johnson, Siaw-Teng Liaw, Marianne Hamilton-Lopez, Daniella Meeker, Toan C. Ong, Patrick Ryan, Ning Shang, Nicole G. Weiskopf, Chunhua Weng, Meredith N. Zozus, and Lisa Schilling. A harmonized data quality assessment terminology and framework for the secondary use of electronic health record data. *eGEMs (Generating Evidence & Methods to improve patient outcomes)*, 4(1):18, 2016.
- [8] John Leider and contributors. Vuetify: Material design component framework. <https://vuetifyjs.com/>.
- [9] OHDSI. OHDSI achilles. <https://github.com/OHDSI/Achilles>, .
- [10] OHDSI. Athena OHDSI vocabulary repository. <https://athena.ohdsi.org/>, .

- [11] OHDSI. Atlas v3.0 source repository. <https://github.com/OHDSI/Atlas3>, .
- [12] OHDSI. Ohdsi atlas (production, atlas 2.x). <https://github.com/OHDSI/Atlas>, .
- [13] OHDSI. OHDSI broadsea. <https://github.com/OHDSI/Broadsea>, .
- [14] OHDSI. OHDSI data quality dashboard (DQD). <https://github.com/OHDSI/DataQualityDashboard>, .
- [15] OHDSI. OHDSI FeatureExtraction r package. <https://ohdsi.github.io/FeatureExtraction/>, .
- [16] OHDSI. OMOP common data model specification. <https://ohdsi.github.io/CommonDataModel/>, .
- [17] OHDSI. OHDSI phenotype library. <https://github.com/OHDSI/PhenotypeLibrary>, .
- [18] OHDSI. PheValuator r package. <https://ohdsi.github.io/PheValuator/>, .
- [19] OHDSI. Ohdsi webapi. <https://github.com/OHDSI/WebAPI>, .
- [20] OHDSI. *The Book of OHDSI: Observational Health Data Sciences and Informatics*. OHDSI Collaborators, 2021. URL <https://ohdsi.github.io/TheBookOfOhdsi/>. Open-access community handbook.
- [21] OHDSI. Ohdsi atlas wiki. <https://github.com/OHDSI/Atlas/wiki>, 2024.
- [22] OHDSI. The book of OHDSI, chapter 16: Clinical validity. <https://ohdsi.github.io/TheBookOfOhdsi/ClinicalValidity.html>, 2024.
- [23] OHDSI. Ohdsi hades (health analytics data-to-evidence suite). <https://ohdsi.github.io/Hades/>, 2024. R packages for cohort generation, characterization, population-level estimation, and patient-level prediction.
- [24] OHDSI. Phenotype Phebruary. Annual community phenotyping sprint, <https://www.ohdsi.org/phenotype-phebruary-2024/>, 2024.
- [25] OHDSI. OHDSI themis: Conventions for the omop common data model. <https://github.com/OHDSI/Themis>, 2024.
- [26] OHDSI. Cohort definition submission requirements — OHDSI phenotype library. <https://ohdsi.github.io/PhenotypeLibrary/articles/CohortDefinitionSubmissionRequirements.html>, 2025.
- [27] OHDSI. Evaluating phenotype algorithms (PheValuator vignette). <https://ohdsi.github.io/PheValuator/articles/EvaluatingPhenotypeAlgorithms.html>, 2025.
- [28] OHDSI HADES. CohortDiagnostics: An r package for performing diagnostics on cohort definitions. <https://ohdsi.github.io/CohortDiagnostics/>, 2025.
- [29] J. Marc Overhage, Patrick B. Ryan, Christian G. Reich, Abraham G. Hartzema, and Paul E. Stang. Validation of a common data model for active safety surveillance

- research. *Journal of the American Medical Informatics Association*, 19(1):54–60, 2012.
- [30] Gowtham A. Rao, Martijn J. Schuemie, Azza Shoaibi, et al. CohortDiagnostics: Phenotype evaluation across a network of observational data sources using population-level characterization. *PLOS ONE*, 2025. doi: 10.1371/journal.pone.0310634.
- [31] William A. Ray. Evaluating medication effects outside of clinical trials: New-user designs. *American Journal of Epidemiology*, 158(9):915–920, 2003.
- [32] Regenstrief Institute. LOINC: Logical observation identifiers names and codes. <https://loinc.org/>.
- [33] Jenna M. Reys, Martijn J. Schuemie, Marc A. Suchard, Patrick B. Ryan, and Peter R. Rijnbeek. Design and implementation of a standardized framework to generate and evaluate patient-level prediction models using observational healthcare data. *Journal of the American Medical Informatics Association*, 25(8):969–975, 2018.
- [34] Martijn J. Schuemie, George Hripcsak, Patrick B. Ryan, David Madigan, and Marc A. Suchard. Empirical confidence interval calibration for population-level effect estimation studies in observational healthcare data. *Proceedings of the National Academy of Sciences*, 115(11):2571–2577, 2018.
- [35] SNOMED International. SNOMED CT. <https://www.snomed.org/>.
- [36] Samy Suissa. Immortal time bias in pharmaco-epidemiology. *American Journal of Epidemiology*, 167(4):492–499, 2008.
- [37] Jill N. Swerdel, George Hripcsak, and Patrick B. Ryan. PheValuator: Development and evaluation of a phenotype algorithm evaluator. *Journal of Biomedical Informatics*, 97:103258, 2019. doi: 10.1016/j.jbi.2019.103258.
- [38] Joel N. Swerdel, Martijn Schuemie, Gayle Murray, and Patrick B. Ryan. PheValuator 2.0: Methodological improvements for the PheValuator approach to semi-automated phenotype algorithm evaluation. *Journal of Biomedical Informatics*, 135: 104177, 2022. doi: 10.1016/j.jbi.2022.104177.
- [39] Joel N. Swerdel, Martijn Schuemie, Gayle Murray, and Patrick B. Ryan. Comparing broad and narrow phenotype algorithms: differences in performance characteristics and immortal time incurred. *JAMIA Open*, 2024. doi: 10.1093/jamiaopen/ooae002.
- [40] U.S. National Library of Medicine. RxNorm. <https://www.nlm.nih.gov/research/umls/rxnorm/>.
- [41] Erica A. Voss, Rupa Makadia, Amy Matcho, Qianli Ma, Christopher Knoll, Martijn Schuemie, Frank J. DeFalco, Ajit Londhe, Vivienne Zhu, and Patrick B. Ryan. Feasibility and utility of applications of the common data model to multiple, dis-

- parate observational health databases. *Journal of the American Medical Informatics Association*, 22(3):553–564, 2015.
- [42] Kazuki Yoshida, Daniel H. Solomon, and Seoyoung C. Kim. The active comparator, new user study design in pharmacoepidemiology: established but no longer the standard? *Current Epidemiology Reports*, 2:221–228, 2015.
- [43] Evan You and contributors. Vue.js: The progressive javascript framework. <https://vuejs.org/>.